

Design Master Plan

Finance Tracker v2/v3 · Investment Intelligence



*“Beauty without logic and logic without beauty
are equally bad.”*

Contents

Part 1	Executive Verdict & Competitive Teardown
Part 2	Principles, Visual Identity, Motion
Part 3	Information Architecture & Page Designs
Part 4	Source UX, AI Interaction, Component System
Part 5	Mobile, Accessibility, Anti-Patterns, Sequencing, QA, Narrative

This bible is planning and specification only. It does not propose code, schemas, or recommendation-logic changes. It is a long-term reference for implementation PRs once Intel v3 is stable and the timing rule in docs/ai/DESIGN_VISION.md is satisfied.

Design Master Plan — Finance Tracker v2/v3

The Quiet Atelier — Design Bible

"Beauty without logic and logic without beauty are equally bad."

This is the long-term design bible for the Investing Intelligence app (Intel + Deploy + Portfolio + Radar + Journal). It is **planning and specification only**. It does not propose code, schemas, or recommendation-logic changes. It is meant to be referenced by future implementation PRs once Intel v3 foundation is stable and the timing rule in `docs/ai/DESIGN_VISION.md` is satisfied.

The bible is split into 5 reference files inside `artifacts/Design_Master_Plan/`:

#	File	Sections covered
00	00_README_and_Verdict.md	Executive verdict, competitive teardown
01	01_Principles_Identity_Motion.md	Design principles, visual identity, motion
02	02_Architecture_and_Pages.md	Information architecture, Command Center, Intel, Deploy, Portfolio, Radar, Journal
03	03_Sources_AI_Components.md	Source/evidence UX, AI interaction, component system
04	04_Mobile_QA_Sequencing_Narrative.md	Mobile, accessibility, anti-patterns, sequencing, QA, north-star narrative

1. Executive Design Verdict

THE DIRECTION HAS A NAME: THE QUIET ATELIER

A **private investment atelier**: the calm of a wealth-management office, the density of an analyst desk, the readability of a luxury editorial publication, and the live pulse of a command center — fused into a single, dark-first interface that respects the reader as a serious thinker, not a retail trader to be entertained.

It is not a dashboard. It is a **morning letter that updates itself**.

WHY THIS BEATS EVERY CATEGORY WE STUDIED

Reference	What they do best	Why The Quiet Atelier wins for <i>this product</i>
Robinhood	Frictionless one-tap action	Confuses entertainment with confidence; gamification erodes trust. We keep the simplicity, drop the dopamine.
Robinhood Legend	Multi-pane charting density	Cluttered, no narrative, no decision guidance. We keep selective density behind progressive disclosure.
Bloomberg Terminal	Information per square inch	Brutalist, hostile to non-experts, 1990s muscle memory. We keep the density discipline and rebuild it as legible editorial.
Monarch / Copilot	Calm typography, soft palettes	Family-budget skew, no analyst depth. We borrow the calm; we add a real intelligence layer.
Wealthfront	Clarity, low-pressure tone	Advisor-paternal, pre-decided for you. We provide guidance without paternalism, with sources.
Public	Storytelling around tickers	Influencer feed, social cosplay. We keep the narrative spine; we cut the social theater.
Seeking Alpha	Depth of evidence	Cluttered UI, paywall walls, opinion roulette. We borrow the source depth and re-author it with composition discipline.
OpenBB / analyst workspaces	Open data, multi-pane analysis	Engineer-skewed, IDE feel. We borrow the data orientation; we throw away the IDE.
Premium editorial (NYT, FT, Monocle, The Browser Co.)	Typographic confidence, hierarchy, restraint	Not transactional or live. We graft live data onto an editorial spine.

Reference	What they do best	Why The Quiet Atelier wins for <i>this</i> product
AI-native (Linear, Arc, Raycast, Notion AI)	Ambient AI, keyboard-first, restrained surfaces	Built for prose/projects, not source-anchored finance. We borrow ambient AI and add source discipline.

WHAT THIS PRODUCT ACTUALLY IS

Internally it is an **investment committee + opportunity radar**. The backend debates, scores, validates sources, challenges itself, and produces deterministic guidance. The frontend's job is the opposite of "more features": its job is to **let one human read the committee's decisions in the time it takes to drink a cup of coffee** — and act with confidence that the math underneath has not been faked.

THE SINGLE DESIGN PROMISE

Within 30 seconds of opening the app, the user should know: 1. What changed in the portfolio overnight. 2. What deserves attention. 3. What action to take (Buy, Hold, Trim, Sell). 4. Why the system thinks that. 5. What evidence supports it. 6. What risks or missing data weaken confidence. 7. How Deploy should act on the next deposit.

Every page, every component, every motion in this bible exists to keep that 30-second promise.

WHAT THE QUIET ATELIER DELIBERATELY IS NOT

- Not a colorful KPI dashboard.
- Not a brokerage trading screen.
- Not a Robinhood clone.
- Not a Bloomberg clone.
- Not a SaaS admin panel.
- Not a social investing feed.
- Not a "famous-investor cosplay" UI.
- Not a chatbot wrapper around stock data.
- Not a generic shadow card layout reskinned with green text.

2. Competitive Teardown

For each reference, in three columns: **Borrow** (good ideas worth stealing), **Reject** (anti-patterns we will not repeat), **Improve** (where we can do meaningfully better).

2.1 ROBINHOOD

Borrow	Reject	Improve
One primary action per ticker view. Clean ticker stream. Big-number-first hierarchy. Calm whitespace on holding pages.	Confetti, gamification, streaks, "rounds up" cuteness. Hidden complexity that erases risk. Loud red/green theatrics. The "Lists" social tab.	Replace one-tap action with one-tap action <i>plus</i> a one-line plain-English thesis. Replace celebratory micro-interactions with quiet confirmation. Replace "Top Movers" with "What changed that matters."

2.2 ROBINHOOD LEGEND

Borrow	Reject	Improve
Multi-pane workspace, dark editorial chrome, dense numerical tables, chart-as-canvas.	Side-panel clutter. Power-user-only ergonomics. Lack of decision guidance. No narrative around the numbers.	Reframe panes as chapters of one document , not floating tools. Each pane must answer a single question. No pane is allowed to exist if you cannot describe its question in one sentence.

2.3 PUBLIC

Borrow	Reject	Improve
Story-around-each-ticker. Calm card density. Treating finance as something narratable.	Influencer feed, "follow this user's portfolio," social cosplay, talking-head videos.	Story is told by the <i>system's own evidence trail</i> , not by personalities. Every narrative claim links to a verifiable source.

2.4 MONARCH MONEY

Borrow	Reject	Improve
Quiet palette, generous typography, calm hierarchy, gentle empty states.	Family-budget mental model. "Spending" rituals. Generic green-blue chart language.	Same calm, but with risk and uncertainty visible instead of softened. Calm should never look like everything is fine.

2.5 COPILOT MONEY

Borrow	Reject	Improve
Tactile micro-motion, gestural drawers, beautifully tuned animations. Soft haptics on iOS.	Slightly playful copy. Consumer-spending skew. Round, illustrative iconography out of place for serious capital.	Same motion vocabulary, drier copy, sharper iconography. Motion explains intelligence , never decorates a flick.

2.6 WEALTHFRONT

Borrow	Reject	Improve
Clarity around fees, allocations, and outcomes. Calm low-pressure tone. Honest projection ranges.	Advisor paternalism ("we've decided for you"). Heavy reliance on illustrative graphics that imply false precision.	Show the <i>system's</i> working: every recommendation has a source trail and a confidence band. The user is co-pilot, not passenger.

2.7 SEEKING ALPHA

Borrow	Reject	Improve
Depth of evidence, willingness to publish a contrarian view, source taxonomy.	Pop-up paywalls, ad density, opinion roulette, no synthesis.	Synthesize multiple sources into one composed thesis. Show contradictions between sources as a feature, not noise.

2.8 BLOOMBERG TERMINAL

Borrow	Reject	Improve
Density per square inch. The discipline that every glyph carries information. The keyboard-first power layer.	Brutalist 1990s aesthetics. Hostile to anyone outside the trading floor. No progressive disclosure. Color-blind unsafe by default.	Keep the density discipline; render it in editorial typography with progressive disclosure. Keep the keyboard layer (■K); rewrite it for plain English.

2.9 OPENBB / ANALYST WORKSPACES

Borrow	Reject	Improve
Multi-source data orientation, willingness to surface raw analyst-grade structure, scriptable workflows.	IDE-skinned UI, dev-only ergonomics, lack of narrative.	Treat the data layer as an <i>editorial source room</i> — analyst-grade behind the scenes, plain English on the page.

2.10 PREMIUM EDITORIAL / LUXURY (NYT, FT WEEKEND, MONOCLE, THE BROWSER COMPANY)

Borrow	Reject	Improve
Typographic system. Engraved rules. Restrained color. Section numerals. Confident whitespace. The reader is treated as intelligent.	Static. No live data. No transactional spine.	Graft live, deterministic data onto the editorial spine. The result is an editorial product <i>that updates itself</i> .

2.11 AI-NATIVE APPS (LINEAR, NOTION AI, ARC, RAYCAST, GRANOLA)

Borrow	Reject	Improve
Ambient AI surfaces. Keyboard-first command bar. AI as utility, not personality. Inline action chips ("explain", "rewrite").	Floating chat bubble. Avatar with a name. Over-reliance on conversation when a structured object is better. AI inventing facts.	AI never invents numbers, scores, allocations, or sources. AI explains, challenges, compares, summarizes — and is labeled when it does.

2.12 CROSS-CUT SYNTHESIS — WHAT WE ARE STEALING FROM EACH, WHAT WE WILL NOT

We are stealing: editorial typography (NYT/FT), dark graphite density (Bloomberg), tactile motion (Copilot), source taxonomy (Seeking Alpha), keyboard-first ambient AI (Linear/Raycast), calm palette (Monarch), and one-decision-per-screen discipline (Robinhood).

We are not building: confetti, streaks, badges, social feeds, illustrated mascots, neon glow, glassmorphism overload, generic shadow cards, advisor paternalism, gamified deposits, raw metric keys in UI, fake price targets, fake confidence percentages, or famous-investor cosplay.

Continued in 01_Principles_Identity_Motion.md.

Design Master Plan — Part 2

Principles, Visual Identity, Motion

3. Product Design Principles

These are hard rules, not taste. Every screen, component, and motion in this bible should be checkable against this list.

- **Trust before excitement.** Visual polish must make the product feel *more* trustworthy, never *more* exciting. If a treatment makes risk look smaller than it is, kill the treatment.
- **One primary decision per screen.** Each screen has exactly one question it is allowed to answer. Secondary information is allowed; secondary primary actions are not.
- **Source-backed claims earn visual weight.** A claim with three reliable sources renders heavier (slightly larger, slightly more saturated) than a claim with one. A claim with no source does not render at all.
- **Missing data is honest, not hidden.** A grey "data missing" pill is always preferable to silent absence or interpolated guesses. The user should never wonder if a number was real.
- **Numbers are sacred.** Tabular figures only. Slashed zeros. Fixed decimal alignment. Currency symbols sized at 0.85× of the digit run. Negative numbers in parentheses *or* with a minus, never both. Every number has a unit. Every number has a timestamp accessible within one click.
- **Motion reveals intelligence, not decoration.** Motion exists to explain what changed, what is loading, or what is about to happen. If a motion has no informational job, remove it.
- **Beginner language, expert structure.** The vocabulary on screen reads as if written for a smart relative who is new to investing. The structure underneath is institutional.
- **Density behind drawers, calm above the fold.** Above the fold is calm and decisive. Density lives in drawers, sheets, and K-summoned panels — earned, not pushed.
- **Never make uncertainty look certain.** Confidence is rendered as a 5-step ladder, never as a percentage above 90 unless backed by a deterministic computation. Probabilistic claims read as probabilistic.
- **Plain English wins, jargon stays backend.** Words like "alpha," "Sharpe," "free cash flow yield" are allowed in tooltips and source drawers, never in the primary copy. Raw metric keys (`pe_ratio_ttm`) never appear in UI.
- **Risk is visible, never buried.** Every page that shows opportunity also shows the risk of acting on that opportunity. Risk gets the same visual hierarchy as opportunity, never less.

- **Confidence is earned, displayed proportionally.** Higher-conviction claims get more visual weight; lower-conviction claims get less. The reader can read confidence at 3 meters distance.
- **Time is a first-class data dimension.** Every claim has a "fresh as of" timestamp. Every recommendation has a "what changed since last time" link. Every chart has a date range visible without hover.
- **The app reads like a private letter.** Tone is direct, considered, not breathless. No exclamation marks in primary copy. No "■" anywhere. No "Congrats!"
- **Beauty must serve logic; logic must inform beauty.** Visual beauty without informational service is decoration. Logic without visual care is brutalism. Both are bugs.

OPERATIONAL COROLLARIES

- **No primary information lives in tooltips.** Tooltips are for vocabulary, not for facts.
- **No KPI tile glow on hover.** Hover surfaces are for affordance, not theater.
- **No animations on numbers that did not change.** A tween implies a delta; a non-delta tween is a lie.
- **No "tap to learn more" without telling the user what they will learn.** Disclosure must be honest about scope.
- **No skeleton loaders that do not match the final content shape.** A skeleton is a *promise* of the layout that follows.

4. Visual Identity System

The system is defined as **tokens**, not bespoke values, so that future implementation can land it in CSS variables and Tailwind config without rework.

4.1 MODE

Dark mode is the canonical mode. The product was designed dark-first because financial data is read in low-light evenings, on phones at night, before the open. Light mode is a deliberate, calm secondary mode — not an afterthought.

DARK MODE — *OBSIDIAN*

A near-black base with a faint green-blue cast. Avoids the "purple-black" SaaS dark and the "blue-black" gaming dark.

Token	Hex (proposed)	Use
bg.canvas	#0A0B0F	Page canvas, behind everything
bg.surface	#10131A	Default card surface

Token	Hex (proposed)	Use
<code>bg.surface.elevated</code>	<code>#161A23</code>	Elevated cards, drawers
<code>bg.surface.high</code>	<code>#1D222C</code>	Modals, command bar
<code>bg.surface.peak</code>	<code>#262C38</code>	Tooltips, top-most overlays
<code>border.subtle</code>	<code>#1F2531</code>	Default 1pt rules
<code>border.default</code>	<code>#2A3140</code>	Inputs, default cards
<code>border.strong</code>	<code>#3A4253</code>	Focus rings, hovered cards

LIGHT MODE — *PAPER*

Warm paper, graphite ink, forest accent. Reads like a private banking letter — not a stark medical screen.

Token	Hex (proposed)	Use
<code>bg.canvas</code>	<code>#FAF7F2</code>	Warm paper
<code>bg.surface</code>	<code>#FFFFFF</code>	Cards
<code>bg.surface.elevated</code>	<code>#F4F1EB</code>	Drawers
<code>border.subtle</code>	<code>#E7E1D6</code>	Rules
<code>border.default</code>	<code>#D8D1C2</code>	Inputs
<code>border.strong</code>	<code>#B8AE9A</code>	Focus, hovered

4.2 COLOR PALETTE

The palette is intentionally narrow. **One signature accent.** Three secondary accents reserved for specific semantic uses. Everything else is grayscale.

SIGNATURE ACCENT — *ATELIER GREEN*

`#2EC27E` (dark mode) / `#138659` (light mode). Used for:

- The single most important affordance on a page.
- The Buy semantic.
- The brand mark.

Atelier Green is **not** used for hover glows, decorative gradients, or "active" states.

SECONDARY ACCENTS

Token	Hex (dark / light)	Reserved for
accent.lapis	#5B7CFF / #3955D6	Information, Hold semantic, AI-composed prose marker
accent.plum	#B47EFF / #7E48D6	Review / "needs human attention"
accent.saffron	#F2A93B / #B97800	Trim / caution
accent.crimson	#D14C5A / #9A2A38	Sell / risk

These accents are reserved. They do not appear decoratively. If you see Plum on the page, it is because something needs human attention — never because plum looked nice.

4.3 SEMANTIC COLORS — BUY / HOLD / TRIM / SELL

The visible action model is **Buy / Hold / Trim / Sell**. Each has a single color, used identically everywhere (badge, border, dot, chart marker, deploy row). Color is **never the only signal** — every semantic is also distinguished by glyph and label.

Action	Dark hex	Light hex	Glyph	Tone
Buy	#2EC27E	#138659	▲ filled	Atelier Green. Earned, not loud.
Hold	#5B7CFF	#3955D6	■ filled	Lapis. Calm, deliberate, "the thesis stands."
Trim	#F2A93B	#B97800	◆ filled	Saffron. Warm, considered — not panic.
Sell	#D14C5A	#9A2A38	▼ filled	Oxblood. Muted, decisive — never Robinhood-red.

*Note: **Sell** deliberately avoids the standard fintech red. Standard red (#FF3B30) reads as alarm. Oxblood crimson reads as decision.*

4.4 RISK AND CONFIDENCE COLORS

Risk is rendered as a 4-tier ladder. Confidence is rendered as a 5-tier ladder. Both are **always paired with a glyph**.

RISK LADDER

Tier	Label	Dark token	Glyph
1	Quiet	grayscale	·
2	Watch	accent.lapis 30%	■
3	Elevated	accent.saffron 60%	■
4	Acute	accent.crimson 80%	■

CONFIDENCE LADDER

Tier	Label	Visualization
1	Sketch	empty ring
2	Tentative	quarter ring
3	Working	half ring
4	Strong	three-quarter ring
5	Settled	full ring

Confidence is **never** rendered as a percentage above 90 unless the underlying score is deterministic and source-backed. The ladder is the canonical render.

4.5 TYPOGRAPHY

The product uses **two type families**, paired like an editorial magazine.

Role	Family (proposal)	Notes
Display / headlines	Tiempos Headline or GT Sectra	A refined, modern serif. Used at 48 / 36 / 28 / 22 px for editorial moments and section openers.
Body / data	Söhne or Inter Tight	A warm geometric sans with strong tabular figures. Used at 16 / 14 / 13 / 11 px.

Role	Family (proposal)	Notes
Numerics	Body family with <code>tabular-nums</code> and <code>slashed-zero</code> opentype features always on	Numbers occupy a fixed grid; never reflow on tween.
Eyebrow / metric label	Body family, 10–11 px, 0.08 em letterspacing, uppercase	Used as section labels, ticker meta, "as of" stamps.

Type tokens (vertical rhythm on a 4-pt baseline):

Token	Size / line-height	Use
<code>display.xl</code>	48 / 56	Cover, north-star headlines
<code>display.lg</code>	36 / 44	Section openers
<code>display.md</code>	28 / 36	Page titles
<code>headline.lg</code>	22 / 30	Card titles
<code>headline.md</code>	18 / 26	Sub-section titles
<code>body.lg</code>	16 / 24	Reading copy
<code>body.md</code>	14 / 22	Default UI
<code>body.sm</code>	13 / 20	Dense tables
<code>caption</code>	11 / 16	Eyebrows, timestamps

4.6 SPACING SCALE

4-pt baseline. Allowed values: **4, 8, 12, 16, 20, 24, 32, 40, 48, 64, 80**. Nothing else exists. If a layout asks for 18 px, the layout is wrong.

4.7 BORDER RADIUS

Token	Radius	Use
<code>radius.sharp</code>	2 px	Data chips, ticker tags, action badges
<code>radius.card</code>	8 px	Cards
<code>radius.panel</code>	12 px	Drawers, multi-pane panels

Token	Radius	Use
<code>radius.modal</code>	20 px	Modals, command bar
<code>radius.pill</code>	999 px	Status pills, filter pills

No half-radii. No 4px, 6px, 10px exceptions.

4.8 ELEVATION AND SHADOW

The product is **almost flat**. Depth is communicated by 1-pt borders and a single low ambient shadow.

Token	Spec	Use
<code>elev.0</code>	none	Page canvas
<code>elev.1</code>	0 1px 2px rgba(0,0,0,0.25) + 1pt border	Default cards
<code>elev.2</code>	0 6px 24px rgba(0,0,0,0.35) + 1pt border	Drawers, modals
<code>elev.3</code>	0 24px 48px rgba(0,0,0,0.45) + 1pt border	Command bar, top-most

No drop-shadow drama. No glowing rings. No multi-layer "neumorphic" trickery.

4.9 GLASS / BLUR

Reserved for **three** surfaces only:

- The top navigation bar (24 px backdrop blur, 70% surface opacity).
- The K command bar (28 px backdrop blur, 60% surface opacity).
- Confirmation overlays for Deploy execution (32 px backdrop blur, 55% surface opacity).

Glass is **not** used for cards. Glass overuse is the single fastest way to make a fintech app look amateur.

4.10 TEXTURE AND BACKGROUND

- **1% monochromatic noise** on hero canvas surfaces (Command Center hero, login). Visually invisible up close; reads as "premium paper" at distance.
- **Faint engraved grid (0.5pt at 4% opacity)** behind charts, optional toggle.
- **No gradients** on serious data surfaces. Gradients are reserved for the brand mark and login moment only.
- **No background images.** Ever.

4.11 CHART STYLING

Charts are **editorial line work**, not "visualization library defaults."

- 1.5 px line weight for the primary series.
- 1 px for secondaries.
- No shadows. No gradients on lines.
- Axis: 0.5 px grid at 6% opacity, only horizontal — never both axes.
- Labels: tabular numerals, caption size, 60% opacity unless hovered.
- Tooltip: a single rule (1 px) drawn vertically through the data, with a small inline panel — never a floating box.
- Color: semantic only (Buy green, Hold lapis, Sell oxblood, neutral graphite). Never rainbow.
- Sparklines: 24 px tall, no axis, no labels, only the line and a single end-dot.
- Negative regions: filled at 8% of stroke color.

4.12 ICONOGRAPHY

- **Single icon family**, 1.5 px stroke, geometric, slightly rationalist (Lucide as base, with selective custom icons for the top-level intents: Intel, Deploy, Radar, Journal).
- **Action icons heavier** (2 px stroke) than navigation icons (1.5 px) — so Buy/Hold/Trim/Sell glyphs read first.
- No filled illustrative icons. No emoji. No mascots.

4.13 EMPTY / LOADING / ERROR STATES

EMPTY — QUIET STATES

A small editorial illustration (line art, single-color, 64 px). Plain copy: *"Nothing to act on yet. The next intelligence run is at 06:30 ET."* Always names what will appear and when.

LOADING — COMPOSING INTELLIGENCE

Skeletons that mimic the final shape of the content (header rule + 3 lines + chip row), with a slow 1.4 s breath. Never generic gray bars. Never a spinner.

For long-running operations (a fresh Run Agents pass), a top-of-page progress affordance: Composing intelligence · 04 of 34 tickers. Plain English, current-of-total format.

ERROR — SPECIFIC, APOLOGETIC, ACTIONABLE

A small inline panel: what happened, what we know, what to try. Never red flood-fill. Never modal block. Never "Something went wrong." Errors always name a follow-up: *"Try Run Agents again, or open the source drawer to see partial data."*

4.14 PREMIUM DETAILS THAT MAKE THE APP FEEL CUSTOM

These are the small touches that distinguish The Quiet Atelier from a generic fintech reskin. None of them is decorative — each carries a small piece of information.

- **Engraved divider rule** between sections — a 0.5 pt line over a 1 pt highlight, hairline elegance.
- **Chapter numerals** at section openers (I, II, III) in display serif — used for Intel chapters, Journal months, Decision History.
- **Tiny build label** in the bottom-right of every page, caption-sized, 30% opacity: as of 06:42 ET · build 412.b.
- **Page-corner timestamp**: top-right of every page in caption: Tuesday · 7 May 2026 · 06:42 ET · NYSE pre-open.
- **A signature mark** on the Command Center hero — a small letterform initial, drawn once per session, that subtly reflects portfolio temperature (one of five expressions: settled, watching, alert, opportunity, deploy-ready). Never face-like, never animated more than once per page load.

5. Motion and Interaction System

Motion exists for one of four reasons:

- To explain **what changed**.
- To indicate **what is loading**.
- To suggest **what is about to happen**.
- To confirm **what just happened**.

Anything else is decoration and is rejected.

5.1 MOTION TOKENS

Token	Curve	Duration	Use
motion.enter	cubic-bezier(0.2, 0.8, 0.2, 1)	240 ms	Cards, drawers entering
motion.exit	cubic-bezier(0.4, 0.0, 1.0, 1.0)	160 ms	Anything leaving
motion.cross	cubic-bezier(0.4, 0.0, 0.2, 1.0)	280 ms	Page transitions
motion.tween	cubic-bezier(0.25, 0.1, 0.25, 1.0)	120 ms	Number tweens

Token	Curve	Duration	Use
<code>motion.emphasis</code>	<code>cubic-bezier(0.34, 1.56, 0.64, 1)</code>	320 ms	Confirmation moments only — used in <5% of motion
<code>motion.breath</code>	sinusoidal	1400 ms loop	Skeletons

5.2 SPECIFIC INTERACTIONS

Surface	Motion	Notes
Page transition	<code>motion.cross</code> : cross-fade with 8 px Y translate	Never slide-in from a side.
Card entry	Stagger 40 ms × n; each card uses <code>motion.enter</code> with 8 → 0 px Y	Stops staggering after card 6 (others arrive instantly).
Thesis update reveal	Characters fade in left-to-right with 4 ms per character; subtle vertical jitter (1 px)	Used only on the <i>single</i> changed thesis on page load — not on every load.
Source drawer reveal	Slide-up 320 ms with <code>motion.enter</code> ; depth shadow appears at the end (last 60 ms)	Drawer always covers ≤ 60% of viewport height.
Risk alert reveal	One-shot pulse on first appearance (scale 1.00 → 1.04 → 1.00, 320 ms emphasis curve); never repeats	The pulse is the alert — the color alone is not enough.
Buy / Hold / Trim / Sell state change	Badge cross-morph in 200 ms; semantic color shift; tabular numerals tween	Optional very-quiet chime, off by default.
Deploy plan confirmation	Ceremonial 600 ms: card lifts 4 px, blueprint hairlines draw from each holding to its ticker target (left → right), then a soft "stamp" appears at the top-right corner	This is the one ceremonial moment in the product. Use sparingly.

Surface	Motion	Notes
Portfolio number change	Numbers tween up/down with <code>motion.tween</code> ; positive change finishes faster (eases out at 100 ms), negative slightly slower (140 ms) — gives losses subtle weight	Never animate a number that did not change.
Hover	1 pt border lift + background shift to <code>surface.elevated</code> ; no scale, no glow	Cards do not bounce.
Skeleton loading	<code>motion.breath</code> at 1.4 s loop; opacity 0.5 ↔ 0.7	No throb, no shimmer-bar.
"New intelligence available"	Top-bar dot fades in (160 ms); side notification slides in from top-right with <code>motion.enter</code> : <i>"3 updates since you were away"</i>	Notification persists 6 s, then collapses to a quiet dot.
Number copying	Selected number gets a 1 pt accent underline that fades 280 ms after copy	Avoids "Copied!" toasts.
Empty → populated	First card uses thesis-update reveal; subsequent cards use card entry stagger	Distinguishes "first arrival" from "subsequent updates."

5.3 REDUCED MOTION MODE

`prefers-reduced-motion: reduce` collapses every motion to a static fade (160 ms, no translate, no scale). Skeletons stop breathing — they show a single quiet shimmer. Deploy ceremony loses its blueprint draw and shows a single "stamped" affirmation at the end.

5.4 MOTION ANTI-PATTERNS BANNED

- No scale-up-on-hover for cards.
- No bounce on landing.
- No confetti.
- No spring physics on data values.
- No parallax.
- No animated charts that "draw themselves" on every render — only on first load and on user-initiated change.

- No skeletons that look like the page is broken.
- No transitions on hover-out (only on hover-in) — out should feel instant.

Continued in 02_Architecture_and_Pages.md.

Design Master Plan — Part 3

Information Architecture and Page Designs

6. Information Architecture

6.1 TOP-LEVEL NAVIGATION

Six primary destinations. One settings/data-health corner. **No "More" menu**, no hidden hamburger items. If something is not in the six, it does not deserve top-level navigation.



Position	Surface	One-line job
1	Today (Command Center)	What changed, what to act on, what to ignore.
2	Intel	The investment-committee view: Buy/Hold/Trim/Sell across watched tickers.
3	Deploy	Allocation of the next deposit, with rationale.
4	Portfolio	Holdings as a living document — concentration, exposure, thesis health.
5	Radar	Opportunity discovery — companies not yet owned.
6	Journal	Decision history and learning loop.
corner	K	Ambient AI + global search + jump.

Position	Surface	One-line job
corner	Data-health dot	One-glance backend / source freshness state.
corner	Settings	Account, sources, brokerage links, preferences.

6.2 NAVIGATION PRINCIPLES

- **Today is always one tap away.** Even from the deepest drawer.
- **Active section reads as engraved**, not highlighted — a 1 pt accent rule under the label, no pill.
- **The data-health dot** turns from *quiet* (default) to *watch* (one stale source) to *acute* (multiple sources stale or backend errors). It is the one persistent risk affordance.
- **No badges with numbers** on top-level nav. Notifications live in the Today rail, not the chrome.

6.3 PAGE-BY-PAGE ARCHITECTURE SUMMARY

For each page below, the structure is described as:

- **Purpose** — the one question this page answers.
- **Primary user question** — the literal sentence the page must answer in the first three seconds.
- **Key modules** — the components that compose the page.
- **Above the fold** — what must be visible without scrolling.
- **Hidden behind drill-down** — what lives in drawers, sheets, modals.
- **Never appears here** — explicit anti-list.
- **Connections** — how this page links to Intel and Deploy.

7. Today — Command Center

The morning brief. Open the app. Read the first paragraph. Decide.

7.1 PURPOSE

Answer the user's morning question in 30 seconds: *what changed, what deserves attention, what is the next best action?*

7.2 PRIMARY USER QUESTION

"Should I do anything different today than yesterday?"

7.3 ABOVE THE FOLD

The above-the-fold layout is a **two-column editorial spread**, not a tile grid.



7.4 KEY MODULES

Module	Job
The Brief	A 3–4 sentence deterministic prose paragraph composed by backend (no LLM). Reads like a private letter.
Act Today	The single highest-confidence Buy / Trim / Sell candidates — capped at 5. Each row is one click into Intel detail.
Risk Pulse	The 1–4 tickers currently in Elevated or Acute risk tier. Always visible if any exist; absent if none.
Deploy Ready	Cash deployable + the system's suggested allocation. Single button to Review Deploy plan.

Module	Job
New Since You Were Away	A small chronological feed of intelligence events since last visit (counts only above the fold; expand below).
Yesterday → Today strip	A bottom-of-fold horizontal strip: portfolio change in dollars, %, and one editorial sentence ("Largest contribution: NVDA +1.8%").

7.5 BELOW THE FOLD

- Full intelligence feed (chronological).
- Source-freshness summary (how many sources fresh, watch, stale).
- Decision Journal preview — last 3 entries.
- Market context strip (S&P, NASDAQ, BTC, 10Y) — caption-sized, never the hero.

7.6 NEVER APPEARS ON TODAY

- Order entry. (This is not a brokerage.)
- Streaks, badges, "you've checked the app 47 days in a row."
- "Top movers" (we curate signal, we do not report the loudest noise).
- KPI tiles. (We have a brief instead.)
- Charts above the fold. (Charts live on Portfolio and ticker drilldowns.)
- Famous-investor quotes.
- News headlines — except when the system itself ingested them as evidence (in which case they appear linked from the Brief, not as a feed).

7.7 CONNECTIONS TO INTEL AND DEPLOY

- Every Act Today row deep-links to that ticker's Intel detail drawer.
- Deploy Ready summarizes Deploy without making decisions on Deploy's behalf.
- The Brief's claims about thesis updates link to the source drawer for the relevant ticker.

8. Intel

The investment committee. What does the system think about each ticker, and why?

8.1 PURPOSE

Render the deterministic Buy / Hold / Trim / Sell verdict for every watched ticker, with thesis, sources, and risk — in a way that respects the existing recommendation engine and posture model already built (`_derive_intel_posture`, `posture_reason`, the BHTS display contract from PR #128 and the simplification PR after it).

8.2 PRIMARY USER QUESTION

"Across everything I follow, what does the system want me to do — and is it right?"

8.3 LAYOUT

A **three-pane workspace**, but each pane reads as a chapter — not as a floating tool.



8.4 FILTER RAIL (LEFT)

- **All / Buy / Hold / Trim / Sell** with counts. **Watchlist** is a separate auxiliary state below the divider — it does not compete with the Buy/Hold/Trim/Sell semantic.
- Sort options: Conviction, Recency of update, Risk tier.
- A **data-health micro panel** at the bottom: count of fresh / watch / stale sources across the Intel set, with one-click drill into source freshness.

8.5 CARD GRID (CENTER)

Each card is the **Action Card** component (defined in `03_Sources_AI_Components.md`):

- Ticker (display serif, 22 px) + small price meta.
- Action chip (Buy / Hold / Trim / Sell) — always one of four, never a posture label.

- Plain-English thesis: max 3 lines.
- Confidence ring (5-step ladder).
- Source freshness dot (green / amber / grey).
- "Why this view?" affordance — a hairline link to the detail drawer.
- "What changed since last time" pill if applicable.

Cards do not bounce on hover. Border lifts to `border.strong`. The selected card's border becomes `accent.lapis`.

8.6 DETAIL DRAWER (RIGHT)

When a card is clicked, a **right-side drawer** slides in occupying 40% of viewport width on desktop, full-screen sheet on mobile.

Drawer sections, in order:

- **Header** — ticker, action chip, confidence ring, "as of" timestamp.
- **Why this view?** — `posture_reason` text, plain English (already produced by `build_posture_reason()`).
- **Plain-English thesis** — composed prose, maximum 6 lines.
- **Evidence trail** — numbered sources with snippet, freshness dot, credibility tier.
- **Risk challenge** — the system's own counter-thesis ("If we are wrong, this is how it breaks").
- **What changed** — diff vs. previous run if applicable; absent if first run.
- **Action footer** — `[Ask why]` `[Challenge]` `[Compare]` chips that summon the AI layer.

8.7 WATCHLIST STATE

Watchlist is **not** a competing label with Buy/Hold/Trim/Sell. It is a separate workflow state for *tickers under research with insufficient evidence to take a side yet*. Visually distinct: dotted border, no action chip, only a "researching" eyebrow.

8.8 NEVER APPEARS ON INTEL

- Raw metric keys (`pe_ratio_ttm`, `_compute_insight_cards` internals).
- Business Read.
- Giant stock report panels.
- Price targets.
- Famous-investor cosplay.
- "Top picks of the week."
- Author bylines, podcast cards, video embeds.

- Social sentiment scores presented as primary truth.

8.9 BACKEND CONTRACT PRESERVED

The visible action model stays exactly **Buy / Hold / Trim / Sell**, normalized through the existing `normalizeDisplayAction` helper. Posture buckets (Add Candidate, Risk Watch, etc.) remain in the data model for backend use, but **never render as filter tabs or card badges**. This bible does not change that contract.

9. Deploy

The execution layer. Convert intelligence into a deposit allocation, with confidence and clarity.

9.1 PURPOSE

Show the user how to allocate the next deposit (or available cash) given the current Intel verdicts. Deploy v2 logic is stable — this is **visual clarity only**.


9.2 PRIMARY USER QUESTION

"My deposit is in. Where does it go, and why?"

9.3 ABOVE THE FOLD

A **single-column ledger** that reads like an investment memo.

[illegible]



9.4 ALLOCATION LEDGER

Each row contains:

- Ticker, action chip (Buy only — Deploy does not deploy Trim/Sell), dollar amount, percent of pool, and a thin allocation bar.
- A one-line rationale connecting back to Intel.
- A link into the relevant Intel detail drawer.

The **Reserve** row is treated as a first-class entry, never as remainder. If cash is reserved, the trigger condition is named.

9.5 CONFIRMATION FLOW

Two-step deliberate flow:

- **Review** — compose a Deploy memo: a single-page summary of all rows with a "what happens next" footer.
- **Confirm** — ceremonial 600 ms motion: blueprint hairlines draw from each Intel verdict to its allocation row, then a "stamped" affirmation appears top-right.

Cancel is always one click away in Review. Confirm cannot fire on hover, double-click, or keypress without focus on the Confirm button.

9.6 DECISION HISTORY SIDEBAR

A right-rail showing the last 5 Deploy executions: date, total amount, rows, and outcome timestamp. Click for full record (jumps to Journal).

9.7 HIDDEN BEHIND DRILL-DOWN

- Per-row allocation math derivation (linked from each row, opens a small sheet).
- Cash-discipline rules in effect.
- Historical reserve triggers and outcomes.

9.8 NEVER APPEARS ON DEPLOY

- Order routing UI. (This is allocation guidance, not order entry.)
- Sliders that let the user "tweak" allocations and override Intel — Deploy presents the deterministic plan; manual override is a separate, intentional act handled via a dedicated "manual override" affordance, never as a primary slider.

- "What if" simulators presenting fake projections.
- Fees-and-taxes calculators dressed as advice.

9.9 LOGIC UNTOUCHED

This bible **does not propose** changes to Deploy's allocation math, reserve logic, or trigger conditions. Visual changes only.

10. Portfolio

The living document of what you own. Reading the portfolio should feel like reading the index of a private collection — calm, honest, and current.

10.1 PURPOSE

Show holdings, concentration, exposure, and thesis health in one calm surface. Not a brokerage trading screen.

10.2 PRIMARY USER QUESTION

"Is my portfolio still consistent with what the system believes?"

10.3 LAYOUT

A two-column layout: holdings ledger on the left, exposure and thesis-health summary on the right.

10.4 HOLDINGS LEDGER (LEFT, 60% WIDTH)

A typeset table — editorial, not spreadsheet:

Ticker	Allocation	Cost basis	Today	Thesis	Source freshness
NVDA	18.4%	\$112,400	+1.8%	Buy ▲ Strong	● fresh
SCHD	12.1%	\$74,200	+0.2%	Buy ▲ Working	● fresh
AAPL	9.8%	\$60,100	-0.3%	Hold ■ Working	● fresh
RIVN	4.2%	\$25,700	-2.1%	Trim ◆ Strong	■ watch
...					

- Tabular nums everywhere.
- Negative numbers in oxblood, with a leading minus, never with parentheses *and* minus.

- Thesis column shows the action chip + confidence tier label.
- Source freshness column shows a single dot — green/amber/grey — never a percentage.

10.5 RIGHT COLUMN

Stacked panels:

- **Concentration** — top 5 holdings as a horizontal bar with names; "concentration risk: Working" label.
- **Exposure by sector** — single horizontal stacked bar with up to 6 sector segments; legend below.
- **Exposure by theme** — same treatment, themes (e.g., "AI infrastructure", "consumer staples", "speculative crypto") as defined backend-side.
- **Thesis health** — single sentence: *"32 of 34 holdings have a Working or Strong thesis. 2 have Tentative thesis health."* Click drills into a sheet listing the tentative ones.
- **Source freshness** — *"96% of holdings have at least one source updated in the last 24 hours."* Click drills into a sheet of stale ones.

10.6 HOLDING DETAIL DRAWER

Click any ticker → right-side drawer:

- Header (ticker, current value, today's change, action chip).
- Thesis health timeline (sparkline, last 30 days, line color = action through time).
- Last 3 decisions on this ticker (linked to Journal).
- Source freshness panel (per-source dot list).
- Stale-thesis warning if applicable: calm caution, *"Thesis hasn't been refreshed since 4 Apr — the next intelligence run will re-evaluate."*

10.7 NEVER APPEARS ON PORTFOLIO

- Buy / Sell order buttons. (No execution here.)
- Real-time chart streaming (we are a tracker, not a terminal).
- Performance attribution graphs that are not source-backed.
- "Beat the market by X%" framing.

10.8 CONNECTIONS

- Every holding row links to its Intel detail drawer.
- Stale-thesis warnings link to the next intelligence run schedule on Today.

11. Radar — Opportunity Discovery

The hunting ground. Companies the user does not yet own, surfaced with discipline.

11.1 PURPOSE

Surface candidate opportunities — companies not currently held — with the same evidence rigor as Intel. Never with the same labels.

11.2 PRIMARY USER QUESTION

"Is there something I should be looking at that I am not already?"

11.3 DISCOVERY FEED

A vertical feed of **Opportunity Cards** — visually distinct from Intel cards (slightly elevated bg, surface.elevated, plum 1 pt accent border, no Buy/Hold/Trim/Sell chip).

Each card carries:

- Ticker + company name (display serif).
- A "Why now" headline (plain English, max 12 words).
- 2-3 source-backed reasons (numbered, each with a freshness dot).
- A "Compare vs. current holdings" affordance that opens a side-by-side sheet.
- Risk red flags if any (calm caution, never alarm).
- A workflow chip: Add to research or Promote to Buy candidate (the latter only available when the system has crossed an evidence threshold).

11.4 WORKFLOW STATES (SEPARATED FROM INTEL LABELS)

Radar uses **workflow states**, not BHTS. This avoids any conflict with the Intel display contract:

State	Visual	Meaning
Surfaced	plum dot	System has surfaced this candidate; no user action yet.
Researching	plum dot, dotted ring	User has added it to research.
Promoted	atelier-green dot	User has promoted it to Buy candidate; now it appears in Intel as a Watchlist item until the next intelligence run upgrades it to Buy/Hold/Trim/Sell.
Dismissed	grey dot	User has dismissed it; remains in archive for 90 days.

These are **workflow states**, never confused with action verdicts.

11.5 "NOT ENOUGH EVIDENCE YET" STATE

If a candidate's evidence threshold is not met for promotion, the card explicitly says so: *"Not enough evidence yet. Two more reliable sources are needed before this can be promoted."* The state is not a failure — it is honest.

11.6 COMPARE DRAWER

When a Radar candidate is opened, the drawer shows a **side-by-side comparison** with the closest current holding (by sector, theme, or correlation as defined backend-side). Three columns: Candidate, Closest holding, Delta. Always source-linked.

11.7 NEVER APPEARS ON RADAR

- "Trending tickers" lists.
- Reddit / Twitter / Discord sentiment.
- Famous-investor portfolio mirroring.
- Fake price targets.
- Crystal-ball language: "could 10x", "next NVDA."

11.8 CONNECTIONS

Promoting a candidate writes a Watchlist entry into the Intel data model. Until the next intelligence run produces a Buy/Hold/Trim/Sell verdict, the ticker shows on Intel only as a Watchlist auxiliary item.

12. Journal — Decision History

The learning loop. What did the system recommend, what did the user do, and what happened next?

12.1 PURPOSE

A long-term record of decisions and their outcomes. The journal makes the user's own track record legible. It also makes the system's own track record legible — including failures.

12.2 PRIMARY USER QUESTION

"Was the system right last time? Was I?"

12.3 LAYOUT

A vertical timeline grouped by month, with chapter numerals (I, II, III...) at each month break.

Tuesday · 7 May 2026

▲ NVDA Buy (\$620) Confirmed
 System: Buy · Strong conviction · 3 sources fresh
 You: Followed
 Outcome: Pending — evaluation window 47 days

◆ RIVN Trim (\$800) Confirmed
 System: Trim · Strong conviction · risk Elevated
 You: Deferred (held position)
 Outcome: Pending — evaluation window 30 days
 Note: thesis weakened, source freshness Watch.

■■■■■■■ chapter II — May 2026 ■■■■■■■■

Friday · 3 May 2026
 ▼ XRP Sell (\$1,200) Confirmed
 ...

12.4 ENTRY ANATOMY

Each Journal entry has six fields:

- **Recommendation** — what the system said (action, conviction, source count).
- **Action taken** — what the user did (Followed / Deferred / Overrode).
- **Outcome** — Pending (with evaluation window), Confirmed (system was right), Mixed, Wrong, or Inconclusive.
- **Thesis health change** — Improved / Stable / Deteriorated.
- **Source freshness change** — Fresh / Stale / Unchanged.
- **Lesson** (optional) — a one-sentence editorial note generated when patterns emerge across multiple entries.

12.5 PENDING EVALUATION WINDOWS

A decision is **never** marked Confirmed/Wrong instantly. Each ticker has a deterministic evaluation window (e.g., 30 days for Trim, 60 days for Buy). Until the window matures, the entry shows Pending — evaluation window N days. **No fake performance metrics.**

12.6 LESSONS SURFACE

A separate panel: "*Lessons learned*" — only renders when the system has identified a recurring pattern (e.g., "Trim recommendations under Acute risk have been right 4 of 5 times in the last 90 days"). Patterns are deterministic, never LLM-generated.

12.7 NEVER APPEARS ON JOURNAL

- "You've made \$X this year!" — performance theater.
- Comparisons to S&P that are not source-cited.
- Streak counters.

- Awards or trophies.
- Ranked leaderboards of any kind.

12.8 CONNECTIONS

Every Journal entry deep-links to:

- The Intel detail drawer at the time of the decision (snapshot).
- The Deploy execution that committed the decision.
- The relevant Portfolio holding's current state.

Continued in 03_Sources_AI_Components.md.

Design Master Plan — Part 4

Source/Evidence UX, AI Interaction, Component System

13. Source and Evidence UX

Source rendering is the moral spine of the product. Every claim that can be sourced must be sourced. Every source must be inspectable. Every source has freshness. Every source has credibility. Contradictions between sources are a feature, not noise.

13.1 INLINE EVIDENCE

Inside any claim text, a source citation appears as a **hairline superscript number** (¹, ², ³) at body weight, accent.lapis. Clicking the superscript opens the source drawer to that source. The superscripts are unobtrusive — readable at distance, ignorable up close.

"NVDA is positioned for continued data-center demand¹, with three fresh sources confirming H100 backlog² and one source noting a softening cycle in consumer GPUs³."

13.2 SOURCE DRAWER

A right-side drawer (consistent with the Intel detail drawer pattern) summoned by clicking any superscript or the "Evidence" affordance on a card.

Each entry in the drawer:

- **Ordinal** (1, 2, 3 ...) matching the inline superscript.
- **Snippet** — 1–3 sentences of the source verbatim, in monospace tabular treatment to signal "raw quotation."
- **Freshness dot** — green (fresh, ≤ 24h), amber (watch, 24–72h), grey (stale, > 72h).
- **Credibility tier** — 5 dots, filled to credibility tier (Tier 1: regulator/issuer filings; Tier 2: paid analyst, well-known; Tier 3: established financial press; Tier 4: independent analyst; Tier 5: aggregated commentary).
- **Source name** with date and URL (URL is not displayed inline; copy-link affordance available).
- **Contradiction tag** if this source disagrees with another in the same evidence trail. Renders as a plum chip: *"Contradicts source 2 on capex outlook."*

13.3 FRESHNESS INDICATORS

Three states, three colors, always paired with text:

Dot	Color	Label	Meaning
●	atelier-green	Fresh	Updated within last 24 hours
■	accent.saffron	Watch	24–72 hours old
■	grey 50%	Stale	> 72 hours old

A stale source does **not** disappear — it persists with its grey dot, so the user knows the system has *not yet* refreshed it, rather than wondering if it ever existed.

13.4 CONTRADICTION DISPLAY

When two sources in the same evidence trail disagree, the drawer renders a **contradiction strip** between them — a thin plum rule with a one-line summary: *"Source 1 (Tier 1, fresh) and Source 4 (Tier 3, watch) disagree on revenue trajectory."* The user sees the disagreement first, decides whose weight to trust.

13.5 "EVIDENCE IS WEAK" DISPLAY

When the total evidence body is below the threshold the system requires for a confident claim, the drawer renders a **calm caution panel**: *"Evidence is currently Tentative — 2 sources, one Tier 4. The next intelligence run will attempt to broaden the evidence base."* Never a red banner, never a warning icon.

13.6 "DATA MISSING" DISPLAY

When a piece of data the user might expect is genuinely missing (e.g., no fresh source for a thesis claim), the slot renders an explicit grey pill: *Data missing — last attempted 3 May.* The pill has a tooltip explaining what was attempted and what would be needed. **Never** an empty -- or zero.

13.7 SOURCE CREDIBILITY DISPLAY

Five-dot ladder. Hovering reveals the tier label. Clicking opens a small sheet defining each tier and what kinds of sources qualify. Credibility tiers are **deterministic**, defined backend-side in a curated registry — the LLM does not invent credibility.

13.8 AVOIDING SOURCE CLUTTER

- Inline superscripts cap at 3 per claim. If a claim has more than 3 supporting sources, the superscripts compress: ¹■■■.
- The drawer paginates after 8 sources.
- Sources are deduplicated across an evidence trail — the same source supporting two claims appears once with two link references.

13.9 SOURCE PROVENANCE FOR AI PROSE

When the system uses a small AI surface to compose prose (e.g., the Brief on Today), the prose itself does not get inline superscripts — instead, a **"Composed from"** affordance below the prose lists the source ordinals used. AI-composed prose is **always** labeled (see §14.6).

14. AI Interaction Design

The AI layer is **ambient, structured, and labeled**. It is never a chatbot, never a personality, never an avatar. It exists to explain, challenge, compare, summarize — never to invent.

14.1 THE NON-NEGOTIABLE CONTRACT

The AI is **forbidden** from inventing or modifying any of the following:

- Numbers (price, percent, allocation, ratio, count).
- Scores or confidence values.
- Source claims or citations.
- Price targets.
- Allocation amounts.
- Risk tier classifications.
- Buy / Hold / Trim / Sell verdicts.

The AI is **permitted** to:

- Explain a thesis or recommendation in plain English, drawing only from the system's own evidence.
- Challenge a thesis (produce a counter-argument), drawing only from the system's own evidence.
- Compare two tickers, drawing only from the system's own data.
- Summarize the brief, the journal, or a ticker's history.
- Translate jargon for beginners.
- Surface the source trail.

Any AI output must include a **"Composed"** mark and a list of source ordinals used. If the AI cannot answer without inventing, it must say so plainly: *"I can't answer this without inventing a number. Try opening the source drawer for the underlying data."*

14.2 WHERE AI APPEARS

The AI surfaces in **four places**, no others:

- **The ■K command bar** — keyboard-first ambient surface for "ask" / "challenge" / "compare" / "explain like I'm a beginner" / "what changed" / "summarize risks" / "find a better opportunity."
- **Inline action chips** at the bottom of the Intel detail drawer: [Ask why] [Challenge] [Compare] [Explain to a beginner].

- **The Brief composition** on Today — composed deterministically (preferred) or via a constrained LLM prompt that takes only system-supplied facts.
- **End-of-day digest** — a small composed summary, source-listed, opt-in.

The AI does **not** appear:

- As a floating chat bubble.
- As a sidebar avatar.
- As a "tutor" overlay.
- On Deploy. (Deploy is deterministic.)
- On Portfolio. (Portfolio is direct.)
- As a typeahead in plain inputs.

14.3 THE ■K COMMAND BAR

Summoned by ■K (or Ctrl+K), or by clicking the small command pill in the top right.



- The bar lives at 28 px backdrop blur with a 60% surface opacity, centered, 720 px wide on desktop, full-width sheet on mobile.
- Recent and Suggestions are populated from the user's history and the current page's context.
- Results render as composed prose with the "**Composed**" mark and source ordinals.
- Long-running queries show a "composing" breath skeleton, never a spinner.

14.4 INLINE ACTION CHIPS

At the bottom of every Intel detail drawer:

[Ask why] [Challenge] [Compare] [Explain to a beginner]

- Ask why — explain the recommendation in plain English.

- **Challenge** — generate a counter-thesis.
- **Compare** — opens a small sheet to pick a comparison ticker.
- **Explain to a beginner** — rewrite the thesis at a 10th-grade reading level.

Each chip outputs into the drawer below, never into a popup.

14.5 THE BRIEF COMPOSITION

Today's Brief paragraph is composed deterministically by the backend whenever possible. When LLM composition is used, it is constrained: the LLM receives only the system's own facts (counts, deltas, ticker names, action verdicts) and is forbidden from adding numerics or claims not in its input. The Brief carries a "Composed" mark and lists the source ordinals.

14.6 THE "COMPOSED" MARK

A subtle iridescent shimmer on the first word of any AI-composed prose, plus a thin caption-sized label below: *Composed · 4 sources used*. The shimmer is a single pass on first render — never a loop. The mark is the user's reliable signal: *"This sentence was synthesized; the underlying data is below."*

14.7 ANTI-PATTERNS BANNED IN AI SURFACE

- **Personality**. The AI does not have a name.
- **Avatars**. No face.
- **Typing animation**. (Mimicking a human typing is dishonest.)
- **"I think..." or "In my opinion..." framings**. The AI does not have opinions; the system has opinions.
- **Emojis in AI output**.
- **Open-ended generation that exceeds the source trail**.
- **Refusing to cite sources**.

15. Component System

The components below are the atoms of the product. Each is defined as: **purpose**, **content hierarchy**, **visual treatment**, **interaction**, and **state matrix**.

15.1 ACTION CARD

The signature card of Intel.

Field	Detail
Purpose	Render one ticker's current Buy / Hold / Trim / Sell verdict with thesis, conviction, and source freshness.

Field	Detail
Content hierarchy	Ticker name (display serif, 22 px) → action chip → 3-line plain-English thesis → confidence ring → source freshness dot → "what changed" pill (conditional).
Visual treatment	bg.surface, 1 pt border.subtle, radius.card, 16 px internal padding. Tabular numerals on the price meta. Action chip uses semantic color (Buy/Hold/Trim/Sell).
Interaction	Hover lifts border to border.strong (no scale). Click opens the detail drawer. Long-press on mobile opens drawer too.
States	Default · Hover · Selected · "What changed" present · Stale source · Insufficient data.

15.2 THESIS CARD

A larger reading surface used inside the Intel detail drawer and on Holding detail.

Field	Detail
Purpose	Render the full plain-English thesis with inline source citations.
Content hierarchy	Display serif headline → composed prose body → "Composed · 4 sources" caption.
Visual treatment	bg.surface.elevated, radius.panel, generous 24 px padding, body.lg type, line-height 1.5.
Interaction	Hover on a superscript reveals a 1-line preview. Click opens source drawer.
States	Default · Insufficient evidence · LLM-composed (with shimmer mark) · Deterministic (no mark).

15.3 SOURCE-BACKED CLAIM

A primitive used inside any composed prose.

Field	Detail
Purpose	Bind a single sentence to one or more sources.
Visual treatment	Inline body text with hairline superscript (1-3 max, or compressed range like ¹ ■■■). The superscript is <code>accent.lapis</code> , 0.85 em, baseline-aligned.
Interaction	Hover reveals 1-line source preview. Click opens drawer to that source.

15.4 RISK CHALLENGE CARD

Field	Detail
Purpose	Render the system's own counter-thesis. <i>"If we are wrong, this is how it breaks."</i>
Content hierarchy	Eyebrow <code>RISK CHALLENGE</code> (caption, plum) → 1–3 sentence challenge prose → linked sources (if any).
Visual treatment	<code>bg.surface</code> , plum 1 pt left rule (4 px), <code>radius.sharp</code> on the rule edge.
Interaction	Click expands to full counter-thesis sheet with sources and risk tier.
States	Default · Acute risk (rule thickens to 6 px crimson) · Elevated · Quiet (omit the card if no challenge).

15.5 OPPORTUNITY CARD (RADAR)

Field	Detail
Purpose	Surface a candidate not yet held.
Content hierarchy	Ticker + name → "Why now" headline (display serif) → 2-3 reasons with freshness dots → workflow chip (Add to research / Promote to Buy candidate).
Visual treatment	<code>bg.surface.elevated</code> , 1 pt plum left border (4 px), <code>radius.card</code> , 20 px padding.

Field	Detail
Interaction	Click opens compare drawer. Workflow chip click triggers state transition.
States	Surfaced · Researching · Promoted · Dismissed · Insufficient evidence.

15.6 DEPLOY INSTRUCTION ROW

Field	Detail
Purpose	One row in the Deploy allocation ledger.
Content hierarchy	Ticker → action chip (Buy only) → dollar amount → percent of pool → allocation bar → 1-line rationale → Intel link.
Visual treatment	Single-row, no card chrome, 1 pt border .subtle rule below. Tabular numerals, fixed decimal alignment.
Interaction	Hover lifts background to bg.surface.elevated. Intel link opens detail drawer.

15.7 CONFIDENCE / CAVEAT BADGE

Field	Detail
Purpose	Render confidence ladder (5 steps) as a compact ring.
Visual treatment	Outer ring 1 pt, inner fill at 0/25/50/75/100% in semantic color of the action. 14 px diameter on cards, 22 px in drawers.
Interaction	Hover reveals tier label ("Strong", "Working", "Tentative", "Sketch", "Settled").

15.8 DATA FRESHNESS INDICATOR

Field	Detail
Purpose	Single-glance freshness signal.

Field	Detail
Visual treatment	6 px filled dot in green/amber/grey. Always paired with text ("Fresh"/"Watch"/"Stale") at caption size when in non-card contexts.

15.9 DECISION TIMELINE

Field	Detail
Purpose	Vertical timeline of decisions on Journal.
Content hierarchy	Date stamp → entry rows → chapter break (monthly).
Visual treatment	A single 1 pt vertical rule on the left. Each entry sits 16 px right of the rule with a 6 px dot at its date stamp.
Interaction	Hover on a dot highlights the row. Click expands the entry.

15.10 INTELLIGENCE FEED ITEM

Field	Detail
Purpose	Single chronological event ("Source freshness drop on AAPL", "New thesis composed for NVDA").
Content hierarchy	Time stamp (caption) → 1-line event sentence → optional link (Intel detail / Source drawer).
Visual treatment	No card chrome. Single line per event.

15.11 COMPARE DRAWER

Field	Detail
Purpose	Side-by-side comparison of two tickers.
Content hierarchy	Two columns: Ticker A
Visual treatment	Bottom sheet on mobile, right drawer on desktop.

15.12 EMPTY / LOADING / ERROR STATES

(Covered in §4.13. Each component has its own variant honoring that pattern.)

15.13 MOBILE CARD

A compact variant of the Action Card for phone widths.

Field	Detail
Visual treatment	Full-width card, 12 px padding, ticker + action chip on row 1, thesis on rows 2-3, confidence ring + freshness dot on row 4.
Interaction	Tap opens bottom sheet (drawer). Long-press opens compare picker.

15.14 TOP NAVIGATION BAR

Field	Detail
Visual treatment	Glass surface (24 px blur, 70% opacity). Brand mark + 6 destination labels + ■K pill + data-health dot + settings cog.
Interaction	Active destination has a 1 pt accent rule under the label. Hovers underline at 50% opacity.

15.15 THE "COMPOSED" MARK

Field	Detail
Purpose	Label any AI-composed prose.
Visual treatment	Iridescent shimmer on the first word (single-pass on render) + caption-sized label below: Composed · N sources used.

Continued in 04_Mobile_QA_Sequencing_Narrative.md.

Design Master Plan — Part 5

Mobile, Accessibility, Anti-Patterns, Sequencing, QA, North-Star Narrative

16. Mobile and Responsive Design

The product is mobile-first by *behavior*, not just by viewport. The phone is where the user actually opens it before market open, on the train, in the kitchen. The phone version must be **as honest, as calm, and as decisive** as the desktop — never "lite."

16.1 BOTTOM NAVIGATION

Four primary destinations as a bottom tab bar (because thumb-reach matters):



- **Today, Intel, Deploy, Portfolio** in the tab bar.
- **Radar** and **Journal** are accessible from Today's secondary rail and from ■K (which on mobile becomes a pull-down search summoned by the title bar).
- No "More" hidden tab. If something's not in the four, it lives one tap inside Today.
- Active tab: 1 pt accent rule above the label, no pill.
- Tab bar uses glass surface (24 px blur, 80% opacity) with 1 pt top border.

16.2 CARD STACK WITH PEEK

On mobile, the Intel page becomes a **vertical card stack** with a 12 px peek of the next card always visible at the bottom. This communicates "more below" without requiring a scrollbar. Cards are full-width with 16 px horizontal padding.

16.3 BUY / HOLD / TRIM / SELL ON PHONE

Action chip moves to the **top-right of the card** (thumb-reach). Tap to open the bottom-sheet detail drawer; the chip itself is informational, not a button — the action is opened via the card body, not the chip alone (avoids accidental drawer summon).

16.4 EVIDENCE DRAWER ON MOBILE

The evidence drawer becomes a **bottom sheet** that occupies up to 80% of viewport height. It supports a swipe-down to dismiss. The sheet has a 4 px handle at the top center. Source list is scrollable inside the sheet.

16.5 DEPLOY CONFIRMATION ON MOBILE

Deploy's two-step (Review → Confirm) is preserved. The Review screen is full-screen, scrollable. The Confirm action requires a deliberate **swipe-to-confirm** affordance (a horizontal slider at the bottom: *"Slide to confirm Deploy of \$1,240"*). Swipe is harder to accidentally trigger than a tap, and feels ceremonial without being gamified.

16.6 AVOIDING CHART CLUTTER

On mobile:

- Charts collapse to **sparklines** by default.
- Full charts are summoned by an explicit "Open chart" affordance in the Holding detail.
- Y-axis labels are dropped; values shown only on hover/tap.
- Multi-series charts collapse to one primary series with a "+N more" summon.

16.7 STICKY "WHAT TO DO TODAY" MINI BAR

On Today, after the user scrolls past the Brief, a small sticky bar appears at the top: *"3 actions today · NVDA Buy, RIVN Trim, SCHD Add."* Tap to scroll back to Act Today, or tap a ticker to open Intel detail directly.

16.8 KEEPING IT ADDICTIVE ON PHONE (WITHOUT GAMIFYING)

Addiction here means *worth opening daily*, not *dopamine loops*. We achieve this by:

- Honest novelty: the Brief paragraph is genuinely different every morning.
- Small surprises in the data, never in the chrome: "Source freshness improved on 3 holdings."
- A morning vs. evening tonal difference: the Brief reads slightly differently after market close (recap mode) vs. before open (preview mode).
- A clear "you are caught up" state when there is nothing to act on: *"You are caught up. The next intelligence run is at 06:30 ET tomorrow."* This is *more* premium than a fake to-do list.

16.9 MOBILE-SPECIFIC ANTI-PATTERNS

- No swipe-to-buy gestures (too easy, too consequential).

- No haptics on success states (we are not a game).
- No bottom-sheet auto-summon (sheets open only on tap).
- No portrait-only lock; landscape works for charts.
- No pull-to-refresh as the **only** way to update — refresh is a manual Run Agents affordance plus the scheduled run; pull-to-refresh is a *secondary* convenience on Today only.

17. Accessibility and Trust

A finance product is a trust product. Accessibility and trust language are not afterthoughts — they are the substrate.

17.1 CONTRAST

- All body text meets **WCAG AAA** (7:1) against its background in both modes.
- All UI text (labels, captions) meets **WCAG AA** (4.5:1) at minimum.
- Disabled states still meet 3:1, with a visible "disabled" affordance (struck-through label, not just dimming).
- Action chips use color *and* glyph *and* label — never color alone.

17.2 COLORBLIND-SAFE SEMANTICS

Every semantic distinction is carried by **at least two channels** — color and shape (glyph, fill style, position). The Buy/Hold/Trim/Sell glyphs (▲, ■, ◆, ▼) are deliberately distinct shapes that read in monochrome. A user with deuteranopia can still distinguish Buy from Trim by glyph, position, and label.

17.3 MOTION REDUCTION

`prefers-reduced-motion:` `reduce` collapses all motion to static fades (160 ms, no translate, no scale). Skeletons stop breathing — they show a single quiet shimmer that fades after 1.4 s instead of looping. The Deploy ceremony's blueprint draw is replaced by an immediate stamped affirmation.

17.4 READABLE FINANCIAL EXPLANATIONS

- Every page has a primary copy block at body.lg (16 / 24) — the "reading copy" of the page. No primary copy below body.md.
- Sentences are kept under 24 words where possible. Multi-clause sentences are split.
- Jargon terms (e.g., "yield curve", "concentration", "drawdown") are **dotted-underlined** on first use, with a tooltip translating to plain English.

17.5 BEGINNER-FRIENDLY COPY

- Action prose is written in second person, present tense: *"NVDA shows three fresh sources supporting a Buy thesis."*
- We never use "we" to mean the system as a personality. The system is referred to as "the system" or "the intelligence layer," never "I" or "we."
- We never use "alpha", "beta", "Sharpe", "drawdown", "Sortino" in primary copy. Backend-defined metrics that have those names render as plain English: "Risk-adjusted return", "downside swing", etc.
- The Explain to a beginner AI affordance exists everywhere a thesis exists.

17.6 REGULATORY AND TRUST CAVEATS

This is a **personal** investing app — not a regulated product. But the user is one human, and their money is real. So:

- The Today Brief footer carries a single quiet caveat: *"This is your private intelligence layer. It is not investment advice. Sources may be wrong. Risks may be larger than displayed."*
- Each Deploy confirmation includes a final-step caveat: *"Confirming will record this decision in your Journal. The system can be wrong. Sources can be stale. Risks can be larger than displayed."*
- No language ever implies certainty. "Will," "guaranteed," "outperform," "beat the market" — banned.

17.7 NO MANIPULATIVE DARK PATTERNS

Banned from this product:

- Urgency theater: "Only 4 minutes left to deploy!"
- False scarcity: "Limited research slots remaining."
- Sticker shock framing: "Your portfolio could have been worth \$X if..."
- Loss-aversion nudges: "Don't miss this opportunity."
- Social proof manipulation: "847 people are watching this stock right now."
- Auto-progression: confirmation flows that auto-advance after a timer.
- Pre-checked consent.
- Hidden cancel buttons.

17.8 NO OVERCONFIDENT LANGUAGE

- The system **never** says "should" without a source. (Compare: *"You should buy NVDA"* vs. *"Three sources support a Buy on NVDA."*)
- The system **never** says "will." (Compare: *"NVDA will rise"* vs. *"NVDA's thesis is currently Strong."*)
- The system **never** says "always" or "never" about market behavior.
- Probabilities, when stated, are stated in deterministic ranges: *"In 12 of the last 15 similar setups, this outcome held."* — never *"There's an 80% chance this works."*

17.9 KEYBOARD NAVIGATION

- Every interactive element is reachable by Tab.
- **■K** opens the command bar.
- **[** and **]** navigate between cards in Intel.
- **?** opens a keyboard-shortcut sheet.
- **Esc** closes any drawer or sheet.
- Focus rings are 2 pt accent.lapis with 2 pt offset, always visible (never out line: none).

17.10 SCREEN READER

- Every action chip carries an `aria-label` that includes the action and the ticker: `aria-label="Buy candidate, NVDA, Strong conviction, 3 fresh sources"`.
- Every freshness dot carries an `aria-label`: `aria-label="Source 2 of 3, fresh, updated 2 hours ago"`.
- Composed prose carries an `aria-describedby` linking to the source drawer count.
- Charts carry an `aria-label` summarizing the trend in plain English.
- Live regions announce: new intelligence available, Deploy confirmation, source freshness changes (these are batched to avoid noise).

18. Anti-Patterns Banned

A strict list. If any future PR introduces any of these, the merge gate must reject.

18.1 VISUAL / AESTHETIC

- Confetti, animated streaks, "level up" badges.
- Neon glow on hover.
- Generic shadow rounded-2xl card with `bg-zinc-900` and a green tag.
- Crypto-casino: gradients, glow, neon-on-black.
- Glassmorphism on cards (reserved for top nav, command bar, deploy confirmation only).
- Heavy drop shadows (multi-layer "neumorphic" trickery).
- Multi-color logos, rainbow charts, *viridis* palettes on serious finance data.
- Stock-photo backgrounds, illustrated mascots, animated mascot reactions.
- Round, illustrative iconography for capital allocation surfaces.
- Animated KPI tile counters everywhere.
- Pulsing dots that loop indefinitely.

- Skeleton loaders that look like bug bars rather than the final layout.

18.2 INFORMATIONAL / HONESTY

- Famous-investor cosplay: "What would Buffett do?", Munger quotes, Lynch invocations.
- Fake price targets.
- Fake confidence percentages (e.g., "82% confidence" when the underlying score is not deterministic).
- Hidden risk under "see more."
- Predictions phrased as guarantees.
- Performance theater: "You've beat the S&P by X%."
- Streak counters, awards, leaderboards.
- Made-up source claims.
- Auto-rounded numbers that hide precision drift.
- Tooltips for primary information.
- Empty -- placeholders where "Data missing" is the truth.
- Exposing raw metric keys (`pe_ratio_ttm`, `intel_filter_bucket`) in UI.
- Conflicting label systems (e.g., showing "Watchlist" as both a Buy/Hold/Trim/Sell action and an auxiliary state simultaneously).
- Business Read panels in UI.
- Giant stock-report blocks.

18.3 TONE / COPY

- Exclamation marks in primary copy.
- Emojis (anywhere outside the user's own free-text journal entries, if those even exist).
- Breathless adjectives: "amazing", "explosive", "huge", "incredible".
- "Pro tip!" / "Did you know?" framings.
- "Congrats!" on any monetary event.
- Second-person imperative phrased as marketing: "Don't miss this!"
- "I" or "we" referring to the system.
- "You should" without a source.

18.4 INTERACTION

- Floating chat bubble.
- Avatar with a name for the AI.

- Auto-progressing confirmation flows.
- Hidden cancel.
- Pre-checked consents.
- Pull-to-refresh as the only way to refresh data.
- Swipe-to-buy gestures.
- Haptic celebration on monetary events.
- Sound effects (other than an opt-in, very-quiet chime on Buy/Hold/Trim/Sell state change, off by default).

18.5 ARCHITECTURAL

- One PR that touches Intel + Deploy + Portfolio + Journal at once.
- A "complete redesign" PR.
- UI changes that touch backend logic.
- New component libraries imported "to support the redesign."
- New animation libraries beyond what's already available (Framer Motion if already in, otherwise CSS transitions).

19. Implementation Sequencing

This is a **proposed** PR sequence for after Intel v3 foundation is stable and the timing rule in `docs/ai/DESIGN_VISION.md` is satisfied. Each PR is small, scoped, reversible. No giant redesign PR.

For each PR: scope, model, risk, UI budget, why this order, how to avoid breaking Intel/Deploy.

#	Scope	Model	Risk	UI budget (max files)	Why this order	Don't break
1	Token + primitive pass. Wire the visual identity tokens (colors, spacing, radii, elevation, motion tokens) into <code>tailwind.config.ts</code> and <code>globals.css</code> . No surface changes yet.	Codex	Low	2 (tailwind + globals)	Foundation. Every later PR depends on these tokens. Cheap, mechanical, reversible.	Existing class names; do not rename existing tokens, only add new ones.
2	Typography pass. Introduce the two type families (display serif + body sans) via CSS variables; apply across global heading and body styles.	Codex	Low	3	Typography sets the editorial spine; everything later inherits.	Existing component class names; opt-in via new utility classes.

#	Scope	Model	Risk	UI budget (max files)	Why this order	Don't break
3	Action Card primitive. Redesign the AgentInsightCard to the spec in §15.1 — same data, new visuals. No backend change.	Sonnet	Medium	2 (card + its test)	Action Card is the most-rendered surface in Intel; refactoring it once unlocks Intel polish.	Backend recommendation pipeline; preserve normalizedDisplayAction contract.
4	Source-Backed Claim primitive + Source drawer. New components, new drawer; wire to existing source data structure.	Sonnet	Medium	3	Source UX is the moral spine; it must exist before further pages depend on it.	The source data contract; this PR consumes only existing fields.
5	Today (Command Center) above the fold. Implement Brief, Act Today, Risk Pulse, Deploy Ready — all consuming existing data.	Sonnet	Medium-High (budget review)	5	Today is the most-opened surface; high-leverage polish.	Deploy logic untouched; Brief composition uses backend-supplied facts.

#	Scope	Model	Risk	UI budget (max files)	Why this order	Don't break
6	Intel detail drawer. Implement the right-side drawer per §8.6, replacing inline expansion.	Sonnet	Medium	4	Drawer pattern unlocks every other detail surface.	The posture / posture_reason contract; consume read-only.
7	Portfolio holdings ledger polish. Convert the holdings table to the editorial spec in §10.4.	Sonnet	Medium	4	Portfolio is the second-most-opened surface; polish after Intel and Today.	Holdings data contract; no new fields.
8	Deploy clarity polish. Apply visual spec to the Deploy ledger (§9.3). Logic untouched.	Sonnet	Medium-High (budget review)	4	Deploy is the most consequential confirmation; polish carefully, after drawer + source patterns proven.	Deploy allocation math; absolutely no logic touch.

#	Scope	Model	Risk	UI budget (max files)	Why this order	Don't break
9	Radar (new page). Discovery feed + Opportunity Card + workflow states. Data contract may need a small backend addition (workflow state enum).	Sonnet	High (split if backend needed)	5	Radar is net-new; lands later because it depends on prior primitives.	Intel display contract — Radar uses workflow states , never Buy/Hold/Trim/Sell labels.
10	Journal (Decision History). Timeline view, entry anatomy, lessons surface.	Sonnet	Medium	4	Net-new view; consumes existing decision log.	Decision log determinism ; no LLM authorship of Confirmed/Wrong outcomes.
11	Motion system pass. Apply motion tokens to entries, transitions, drawers. Reduced-motion respected.	Sonnet	Medium	3	Motion is layered last so prior PRs are validated static-first.	Performance; motion must not regress paint times.

#	Scope	Model	Risk	UI budget (max files)	Why this order	Don't break
12	Mobile pass. Bottom tab bar, card stack, bottom-sheet drawers, swipe-to-confirm.	Sonnet	Medium-High (budget review)	5	Mobile is layered after desktop polish so the design is canonical first.	Existing routes; mobile is layout-only.
13	AI ambient layer (■K + chips). Implement command bar + inline action chips with the AI contract from §14.1.	Sonnet	Medium	4	AI sits on top; lands when source UX is solid so AI can label its citations.	The AI contract — never invents numbers; LLM prompts must be source-bounded.
14	Empty / loading / error pass. Replace generic states with the "Quiet states" / "Composing intelligence" spec.	Codex	Low	4	Polish; lands late because each surface has its own empty state.	Existing fallback behaviors; no logic change.
15	Iconography swap. Apply the bespoke top-level intent icons + heavier action icons.	Codex	Low	3	Final aesthetic detail.	Accessibility; ensure aria-labels remain.

After each PR: **Codex cheap visual merge gate** per `docs/ai/skills/ui_fix.md`.

After any Medium-High or High PR: **stop the Sonnet session**; bring the PR back to ChatGPT/Codex for review.

If at any point a PR scope grows beyond its `Max files cap`, split.

19.1 WHY NOT ONE GIANT REDESIGN PR

- A giant redesign PR is unreviewable.
- It will conflict with any in-flight Intel/Deploy work.
- It will hide regressions in mass diff.
- It will exhaust Sonnet usage in one shot.
- It contradicts the project's `docs/ai/PROMPT_LIBRARY.md` UI budget gate.

19.2 PRE-CONDITIONS

Before PR #1 starts:

- Intel v3 recommendation engine is stable (no churn in `recommendation_engine.py` for at least 2 weeks).
- Deploy v2 is stable (no churn in allocation math).
- Decision log persistence is stable.
- The timing rule in `docs/ai/DESIGN_VISION.md` is satisfied.
- The user has reviewed and approved this bible (this artifact).

20. Design QA Checklist

Every UI PR before merge must pass this checklist. The merge gate (Codex per `docs/ai/skills/ui_fix.md`) reads this list against the diff.

20.1 VISUAL QUALITY

- [] All colors come from the token palette — no raw hex.
- [] All spacing comes from the 4-pt scale — no off-grid values.
- [] All radii come from the 5 defined tokens.
- [] All elevation comes from the 4 defined tokens.
- [] Glass/blur used only on top nav, command bar, or Deploy confirmation.
- [] No drop shadows beyond `elev.0-elev.3`.
- [] Charts use semantic color only, no rainbow.

- [] Sparklines have no axis or labels.
- [] Tabular numerals on every numeric column.

20.2 DATA CORRECTNESS

- [] Every claim is source-backed or marked "Data missing."
- [] Every number has a unit and a timestamp accessible within one click.
- [] No -- placeholders.
- [] No interpolated/guessed values.
- [] No rounding that hides precision drift on monetary values.

20.3 SOURCE CORRECTNESS

- [] Inline superscripts cap at 3 or compress (^{1■■■}).
- [] Every superscript opens the source drawer to that source.
- [] Source drawer shows snippet, freshness, credibility, contradictions.
- [] Stale sources persist with grey dot, never disappear.

20.4 NO RAW METRIC KEYS

- [] No `pe_ratio_ttm`, `_compute_insight_cards`, `intel_filter_bucket`, etc., visible in UI.
- [] Backend names translated to plain English.

20.5 NO CONFLICTING LABELS

- [] Buy / Hold / Trim / Sell are the only visible action verdicts.
- [] Watchlist is rendered as auxiliary state, visually distinct.
- [] Radar uses workflow states (Surfaced / Researching / Promoted / Dismissed), not BHTS.
- [] Posture buckets do not appear on cards or filters.

20.6 NO BROKEN DEPLOY LOGIC

- [] No changes to `recommendation_engine.py` allocation math.
- [] No changes to reserve trigger conditions.
- [] Deploy logic is consumed read-only by the UI.

20.7 RESPONSIVE BEHAVIOR

- [] Mobile bottom nav has 4 entries; no hidden "More."
- [] Drawers become bottom sheets on mobile.

- ☐ Confirm flows use swipe-to-confirm on mobile.
- ☐ Charts collapse to sparklines.
- ☐ No horizontal scroll on any default page width ≥ 320 px.

20.8 ACCESSIBILITY

- ☐ Body text meets WCAG AAA (7:1).
- ☐ UI text meets WCAG AA (4.5:1).
- ☐ Every action chip has aria-label including action, ticker, conviction.
- ☐ Every freshness dot has aria-label.
- ☐ Charts have aria-label summarizing trend.
- ☐ `prefers-reduced-motion`: reduce collapses all motion.
- ☐ Focus rings visible on every interactive element.
- ☐ Keyboard navigation reachable for every action.

20.9 LOADING / ERROR STATES

- ☐ No spinners.
- ☐ Skeletons mimic final content shape.
- ☐ Errors specific, apologetic, actionable.
- ☐ Long-running ops show "Composing intelligence · N of M."

20.10 PLAIN-ENGLISH COPY

- ☐ No "alpha", "Sharpe", "drawdown" in primary copy.
- ☐ Sentences ≤ 24 words where possible.
- ☐ Jargon dotted-underlined with tooltip translation.
- ☐ No exclamation marks in primary copy.
- ☐ No emojis in UI.
- ☐ No "you should" without a source.
- ☐ No "will" / "guaranteed" / "outperform" language.

20.11 DARK / LIGHT MODE

- ☐ All tokens defined in both modes.
- ☐ Semantic colors readable in both modes.
- ☐ No hard-coded `#fff` or `#000` in components.

20.12 AI LABELING

- [] All AI-composed prose carries the "Composed" mark.
- [] All AI-composed prose lists source ordinals used.
- [] No AI invents numbers, scores, allocations, sources, or verdicts.

20.13 ANTI-PATTERN ABSENCE

- [] No confetti, streaks, leaderboards, badges.
- [] No floating chat bubble.
- [] No avatars.
- [] No fake price targets.
- [] No fake confidence percentages.
- [] No famous-investor cosplay.
- [] No hidden risk.
- [] No urgency theater.

21. Final North-Star Narrative

*Tuesday morning. 6:42 ET. Pre-open. I open the app on my phone. Glass top bar; the brand mark sits on the left, a small letterform initial drawn once for this session. The data-health dot is quiet green. The clock and the build label are tucked in the corner — calm, never decorative. The page is titled simply **THE BRIEF**, in the display serif. Below it, three editorial sentences: Three thesis updates since yesterday's close. One acute risk on RIVN. The portfolio is up 0.4% in pre-market on broad strength. Cash to deploy: \$1,240. The next intelligence run completes at 07:00 ET. The shimmer on the first word is the gentle "Composed" mark. Below the prose, in caption: Composed · 7 sources used. An engraved divider rule, then two columns. On the left, **Act Today**. Three rows. NVDA — Buy, three sources, fresh. SCHD — Add to position. RIVN — Trim, risk Elevated. Each row is one tap into Intel detail. On the right, **Risk Pulse**. RIVN at Acute. BTC and KLAR at Elevated. The acute row sits at the top with a thin oxblood left rule — the only oxblood on the page. I tap RIVN. The right edge of my screen slides up a bottom sheet — the Intel detail drawer, mobile form. Header: ticker, action chip (Trim, saffron), confidence ring three-quarters filled. Why this view? — "Six fresh sources have shifted from Watch to Stale on Q1 delivery numbers. The remaining fresh source notes a softening order book." Risk challenge: a plum left rule. "If we are wrong, this is how it breaks: the order book recovers on the consumer EV cycle, which the data does not yet support." What changed: a small chip — "Conviction moved from Working to Strong since 5 May." I tap it; a tiny diff sheet shows the previous thesis next to the current one. I tap [Ask*

why]. The drawer extends with a composed paragraph, the Composed mark on its first word, three source ordinals listed below. No personality. No avatar. Just the system, explaining itself. I close the drawer. Back on Today, I scroll to **New since you were away**. Three thesis updates. One new Radar candidate. I open Radar. A vertical feed of candidate cards on bg.surface.elevated with thin plum left borders. The new candidate is a small specialty chip designer I had never heard of. Why now: "Three Tier-2 sources cite a contract with a Tier-1 cloud provider; the closest comparable holding in your portfolio is NVDA at 18.4%." I tap **Compare vs. current holdings**. A side-by-side sheet. A delta column. Honest. I tap **Add to research**. The card's workflow chip moves from Surfaced to Researching. No celebration. No confetti. Just a quiet, decisive state change. I switch to **Deploy**. \$1,240 deployable. The page reads as a single-column ledger. NVDA — Buy — \$620 — 50% — "Three fresh sources support thesis. Conviction Strong." SCHD — Buy — \$400 — 32% — "DCA target. Yield discipline preserved." Reserve — \$220 — 18% — "Held for opportunity. Trigger: any Acute risk drops to Elevated." I tap **Review**. A full-screen memo. I tap-and-hold the swipe-to-confirm slider at the bottom: "Slide to confirm Deploy of \$1,240." The slider eases under my thumb; at the end, the ceremonial 600 ms motion plays — the blueprint hairlines draw from each Intel verdict to its allocation row, a soft stamp appears top-right. The page settles. A new entry appears in my **Journal**. Pending — evaluation window 47 days. I close the app. The whole interaction took 94 seconds. I have read what changed, understood what mattered, seen the evidence, considered the risk, made one deliberate decision, and recorded it for the future me to learn from. The system was there as a quiet, source-anchored co-pilot. It did not entertain me. It did not pretend to know more than it did. It did not invent numbers. It did not show me what other people were doing. It looked beautiful, because every pixel had a purpose. It was useful, because every claim had a source. It was trusted, because uncertainty was visible and honest. It was mine, because the journal will remember what I did and what came of it — better than I will. **This is The Quiet Atelier.**

End of Design Master Plan.