

Finance Tracker Intel v3

Revised Architecture and Build Plan - Draft 3

Advisor revision of Opus Draft 2 for prashanthkrishnan91/claude_financetracker_pk91

Date: 2026-05-04 | Stack: Next.js + FastAPI + Supabase | Scope: v2 only unless a future architecture phase explicitly says otherwise

Executive decision: Keep the v3 thesis, but replace the 19-PR hedge-fund-terminal rollout with a kernel-first implementation path. The first implementation PR is not Data Truth, not EDGAR, and not a forest of empty stubs. It is the action-contract lock plus HOLD-collapse canary.

This PDF supersedes Draft 2 as the build reference. Draft 2 remains useful as a vision document. Draft 3 is the execution contract for prompts, model routing, PR sequence, budget gates, and stop rules.

Core product promise: A personal investment committee and opportunity radar that improves decision quality per unit of evidence. It does not promise alpha, certainty, price targets, or automated trading.

Immediate next move: Codex, new chat, Low usage: PR 0 - lock visible Intel actions to BUY / HOLD / TRIM / SELL, block posture-label regressions, add HOLD-collapse canary, update HANDOFF and progress log, stop after PR.

Contents

- 1. Executive verdict and correction to Draft 2
- 2. Non-negotiable product contract
- 3. What Draft 2 got right, wrong, and missing
- 4. Revised architecture: v3 kernel before v3 universe
- 5. Deterministic action policy contract
- 6. Data, sources, snapshots, and citations
- 7. LLM committee: explain and challenge, never decide
- 8. Opportunity Radar: keep, but do not start there
- 9. Frontend product contract
- 10. Deploy boundary
- 11. Correct PR sequence
- 12. Prompt and model routing rules
- 13. Immediate PR 0 prompt
- 14. Phase gates and acceptance criteria
- 15. Module backlog: must, later, reject
- 16. Risks, stop rules, and recovery plan
- 17. Final recommendation

1. Executive verdict and correction to Draft 2

Draft 2 correctly identifies the fundamental product failure: Intel v2 became a posture-patching loop. It produced label churn, page-load/run inconsistency, and eventually the exact bad canary state: all cards showing HOLD again. The fix is not another renamed posture label. The fix is a deterministic, source-backed action engine with clean frontend projection.

Verdict: Accept Draft 2 as a vision document. Reject Draft 2 as the literal implementation sequence. The architecture needs a smaller v3 kernel first, then source packs, then committee, then radar, then Deploy handoff.

The big correction

Draft 2 says start with UI contract lock plus v3 scaffolding, then jump into Data Truth and Snapshot Store. That is still too heavy and too noisy for the first move. Draft 3 changes the first implementation step to a pure guardrail PR:

- Lock the visible action model.
- Block old posture labels from reappearing.
- Add a HOLD-collapse canary with synthetic data.
- Do not create dozens of empty module files.
- Do not touch Deploy.
- Do not add SQL.

Why this matters

The failed pattern has been broad patching against symptoms. If PR 0 starts by creating 30-40 stub files and touching both frontend and backend loosely, the repo gets more complexity without solving the user-visible contract. The first PR must make the next regression harder, not make the architecture look bigger.

2. Non-negotiable product contract

These rules are the product constitution for Intel v3. Future prompts should quote this section when scope is ambiguous.

- **Intel decides what deserves capital and why.** Deploy decides how much and when. Intel never owns allocation math. Deploy never invents thesis text.
- **Visible held actions are locked:** BUY, HOLD, TRIM, SELL. No Add Candidate, Review, Watchlist, Risk Watch, Trim Candidate, Strong Buy, or Strong Sell.
- **Non-held opportunity actions are later-only:** WATCH, BUY, AVOID. They appear only inside Opportunity Radar after closed-universe validation exists.
- **Numbers are deterministic.** LLMs cannot invent metrics, scores, source claims, price targets, allocations, or tickers.
- **Claims require sources.** A claim without provenance is hidden, deferred, or marked as missing.
- **Missing data is honest.** Suppress an axis; do not collapse every card into identical HOLD/LOW.
- **Advanced metric names stay backend-only.** UI uses plain English, not raw keys or analyst jargon.
- **No giant stock report UI.** The product is a decision cockpit with progressive disclosure, not a 30-section report page.
- **No famous-investor cosplay.** 13F and insider data can be evidence, never a headline like Buffett would buy this.
- **No fake precision.** Confidence is LOW / MEDIUM / HIGH. No percentages, no guaranteed predictions, no precision price targets.
- **Every meaningful PR updates docs/ai/HANDOFF.md and v2/progress_log.md.** README.md only for public/setup/deployment/architecture changes.

3. What Draft 2 got right, wrong, and missing

Category	Draft 2 Item	Draft 3 Decision
Keep	Source-backed deterministic action policy	Core architecture. This is the heart of v3.
Keep	LLM committee cannot emit action labels	Non-negotiable. LLMs explain and debate only.
Keep	Per-axis suppression instead of global insufficient_data	Required to prevent HOLD-34.
Keep	Closed-universe Opportunity Radar	High-value differentiator, but later than Draft 2 suggests.
Revise	19 PR roadmap	Replace with 12-13 tighter PRs and phase gates.
Revise	PR 0 creates dozens of stub files	Reject. Use contract lock and canary only.
Revise	Five modules per backend PR	Split. Modules must be grouped by risk and action value.
Defer	Full EDGAR forensic suite early	Valuable, but heavy. Ship after decision kernel and snapshot read path.
Defer	Replay/backtesting/admin tooling	Phase 2+. Nice discipline, not needed before v3 cards work.
Reject	Any new visible posture buckets	Would recreate Intel v2 inconsistency.

Missing from Draft 2

- **Minimum viable v3 kernel.** Draft 2 jumps from vision to many modules. It needs a tiny decision engine first.
- **Feature flag discipline.** New decision policy must dark-launch behind an explicit flag before replacing cards.
- **Synthetic fixtures.** Tests cannot rely on live market/provider behavior.
- **Module admission rule.** A module is allowed only if it feeds an action axis, veto, source pack, UI sentence, or Deploy handoff field.
- **Explicit anti-scaffold rule.** Empty files are not architecture. Build modules only when their contract and tests ship in the same PR.

4. Revised architecture: v3 kernel before v3 universe

Draft 2 uses a 9-layer architecture from Investment Policy Statement through Replay. Draft 3 keeps the conceptual layers but changes the execution order. The app needs a working decision spine before it needs every professional research feature.

Draft 3 Layer	Purpose	First shipped in
K0 Action Contract	Lock visible actions and block posture regressions	PR 0
K1 Decision Kernel	Emit deterministic action, conviction, blockers, why/why_not_now from existing data	PR 1
K2 Portfolio Fit Lite	Use current weight, concentration cap, under/overweight, role in portfolio	PR 2
K3 Snapshot Lite	Persist immutable v3 decisions and module outputs	PR 3
K4 Held Card UI v3	Show simple action cards and drawer from v3 projection	PR 4
K5 Source Pack + Validator	Build citation grammar and reject hallucinated claims before LLM rendering	PR 5
K6 Committee Lite	Bull, Bear, Risk, PM note only; no action labels	PR 6

Draft 3 Layer	Purpose	First shipped in
K7 Intelligence Modules	Quality, valuation, technical, then forensic modules	PR 7-8
K8 Opportunity Radar	Closed universe and screens after held-ticker engine is stable	PR 9-10
K9 Deploy Handoff	Typed handoff only after Intel v3 output is stable	PR 11
K10 Feedback	Decision logging and performance windows after decisions are stable	PR 12+

Architecture rule: No layer may skip the layer below it. Example: no LLM committee UI before source-pack validators; no Opportunity Radar before closed universe; no Deploy integration before v3 decisions are stable.

5. Deterministic action policy contract

This is the core of v3. The action label must come from a deterministic policy, never from a prompt. The policy should operate on ordinal bands, not a single composite score. This preserves differentiation when one source is missing.

Held ticker action set

Action	Meaning	Required principle
BUY	Add capital if Deploy has room and timing allows	Evidence and attractiveness are at least OK, price is fair/cheap or explicitly explainable, fit is not blocked, no critical veto.
HOLD	Keep position; no new capital priority	Default state. Not a failure. Must still carry useful reason and what-would-change text.
TRIM	Reduce exposure or avoid adding more	Triggered by overweight, valuation/fullness, high risk flag, thesis drift, or better fit elsewhere.
SELL	Exit candidate	Requires sourced broken thesis, serious filing/dilution/earnings-quality veto, or explicit user policy condition.

Primary axes

Axis	Band examples	Purpose
Evidence quality	thin / ok / strong	Determines whether action can be confident. Thin forbids BUY.
Attractiveness	weak / ok / strong	Business quality and return-driver appeal.
Price context	expensive / full / fair / cheap / suppressed	Prevents quality-at-any-price errors.
Portfolio fit	blocked / underweight / on_target / overweight / breach	Turns generic stock view into personal portfolio decision.
Risk veto	none / low / medium / high / critical	Caps or blocks actions when sourced flags exist.
Technical/timing	confirms / neutral / opposes	A modifier only. Never primary.
Thesis status	missing / intact / drifting / broken	Tracks whether the original reason to own still holds.

Priority order

The policy should evaluate actions in this order: SELL, then TRIM, then BUY, then HOLD. This prevents a strong business-quality signal from masking a sourced veto or concentration breach.

- SELL can trigger only on sourced severe conditions or explicit policy.

- TRIM can trigger on concentration, expensive price with non-strong attractiveness, thesis drift, or high sourced risk.
- BUY requires enough evidence, acceptable price, non-blocked fit, and no high/critical veto.
- HOLD is the safe default and must still explain what is missing or what would change the view.

HOLD-collapse canary

Regression rule: A synthetic non-degenerate recommendation set with at least one meaningful published signal must not collapse to all HOLD with LOW conviction. If it does, CI should fail. This protects against the exact HOLD-34 failure class without depending on live portfolio data.

6. Data, sources, snapshots, and citations

Draft 2 is right that the final system needs a Data Truth Layer and immutable snapshots. Draft 3 changes when and how it lands: start with a lite v3 decision snapshot after the action policy is proven, then expand to per-fact source rows and provider health.

Data Truth Lite first

- Do not require every provider and every fact schema before v3 cards can improve.
- Start by wrapping existing facts already used by Intel v2 and adding provenance where available.
- Every new fact should eventually carry value, unit, source, source_url, observed timestamp, published timestamp, freshness, and confidence.
- NaN, None, inf, stale, missing, and provider errors must produce suppressed facts, not fake zeros.
- Provider conflicts should be logged and resolved by source priority, not hand-waved in UI.

Snapshot Store Lite

- First version persists v3 decision output and axis/module payloads by run_id and ticker.
- Snapshots are immutable after write; new refresh creates a new snapshot.
- Cards render from one snapshot id, not blended fallback rows.
- Every snapshot carries schema_version and partial flag.
- Legacy v2 reads can stay for one cycle behind a feature flag.

SQL sequencing

Migration	When	Purpose
016_intel_v3_decision_snapshots.sql	PR 3	Decision snapshot and module output rows. Additive only.
017_policy_versions.sql	After held cards work or if governor needs persisted policy	User investment policy defaults and versioning.
018_decisions_v3.sql	After v3 cards stable	Audit trail and performance windows.
019_theses.sql	Later	Thesis tracker and drift detection.

7. LLM committee: explain and challenge, never decide

The committee is valuable, but only after source-pack validation exists. Draft 2 specifies nine personas. Draft 3 starts with four. More personas do not improve trust if the validator is weak.

Committee Lite

Persona	Allowed	Forbidden
Bull	Strongest ownership case from source pack facts and module outputs	Action labels, unsourced metrics, price targets
Bear	Strongest case against ownership; must be non-empty	Generic risks, no-risks language, invented flags
Risk Officer	Sourced vetoes and material risks only	Speculation without source
PM Note	Plain-English synthesis, why_now, why_not_now, what_would_change	Changing the deterministic action

Validator requirements

- Citation resolution: every citation id exists in source pack.
- Numeric grounding: every number in text must match a source-pack value within tolerance.
- Forbidden phrases: no strong-buy, load up, guaranteed, price target, will go to, or raw metric keys.
- No action overwrite: committee output schema does not contain action.
- Bear non-empty: a missing bear case is rejected and retried or marked deferred.
- Cost cap: if the token budget is exceeded, committee is deferred and deterministic decisions still render.

Do not invert the order. Source pack first, validator second, LLM third, UI rendering fourth. If the prompt runs before the source pack is validated, v3 becomes another professional-looking hallucination engine.

8. Opportunity Radar: keep, but do not start there

Opportunity Radar is the differentiator. Without it, the app only reacts to the current 34 holdings. With it, the app can surface hidden gems and better-fit replacements. But it is a later phase because it multiplies cost, source needs, and hallucination risk.

Radar contract

- Closed universe only. LLM cannot invent tickers.
- Screens are deterministic. LLM can explain screened candidates, not create them.
- Held tickers are excluded or marked as already owned.
- WATCH is the default surfaced radar action. BUY requires standalone action criteria and portfolio fit.
- AVOID appears only with sourced disqualifying flags and is hidden by default behind Show avoided.
- Initial universe should be small enough to inspect: held ETF constituents, manual watchlist, S&P; 500/major sector leaders, and curated quality list. Do not start at 3000 tickers.

First radar screens

Screen	Why it matters	When
Quality compounder	Find durable businesses outside current book	PR 9
Valuation dislocation	Avoid paying too much for quality	PR 9
Momentum with fundamentals	Avoid value traps and meme-only moves	PR 9
ETF upgrade	Find lower-cost/lower-overlap alternatives	PR 10+
Insider/13F clustering	Useful evidence, not headline	Later
Hidden gem small cap	Interesting but false-positive heavy	Later

9. Frontend product contract

The frontend should be a premium decision cockpit, not a report dump. Advanced finance terms can exist in backend payloads, but the UI translates them into beginner-friendly plain English.

Phase A: stable held cards first

- Header: snapshot as-of timestamp, modules ready count, data quality banner.
- Filter tabs: ALL / BUY / HOLD / TRIM / SELL only.
- Each card: ticker/name, action badge, conviction dots, one-sentence why, one-sentence risk, fit chip, evidence chip.
- Drawer: why, risk, what would change the view, how it fits my portfolio, evidence check, what changed.

Phase B: committee drawer

- Bull and Bear side by side.
- Risk Officer bullets with citations.
- PM note in plain English.
- Committee deferred badge when cost cap or missing data suppresses LLM output.

Phase C: radar

- Opportunity Radar section below Current Holdings.
- Rows max 5 by default with show more.
- WATCH / BUY only by default; AVOID hidden behind toggle.
- From screen tag is plain English: quality compounder, better valuation, ETF upgrade. Do not show backend module names.

UI must never show

- Raw metric keys
- Business Read
- backend module names as UI labels
- posture buckets
- famous investor headlines
- price targets

- confidence percentages
- score numbers as the main product
- social feed noise
- act now CTAs

10. Deploy boundary

Deploy v2 is stable. Do not touch it while Intel v3 is still learning to emit reliable actions. The boundary is a typed handoff, added only after v3 decisions and frontend rendering are stable.

IntelToDeploy handoff - later

Field group	Examples	Rule
Ticker decision	ticker, action, conviction_band, why_now	Produced by Intel decision policy only.
Capital priority	urgent_buy, opportunistic_buy, hold, trim_first, sell_first	Produced by portfolio governor, not UI.
Risk limits	max_position_pct, must_not_exceed_pct, blockers	Deploy consumes, does not invent.
Suppressions	missing valuation, thin evidence, stale data	Deploy can choose to avoid capital on thin names.
Portfolio context	cash_pct, risk_state, recent_drawdown	Used by Deploy timing/math only.

Stop rule: If an implementation prompt touches Deploy before PR 11 or before v3 held-card output is stable, reject or split it.

11. Correct PR sequence

This sequence is optimized for one-pass correctness, token efficiency, and avoiding another patch loop. Each PR is one chat, one model choice, one branch/PR, then stop.

P R	Scope	Model	Usage	UI	SQL	Stop gate
0	Action contract lock + posture-label block + HOLD-collapse canary	Codex	Low	Tiny	No	Filters/badges locked; no stubs; docs updated
1	Decision Kernel Lite dark launch using existing data	Sonnet	Medium	No	No	BUY/HOLD/TRIM/SELL fixture matrix passes
2	Portfolio Governor Lite: concentration and fit bands	Sonnet	Medium	No	No	Fit affects action without Deploy changes
3	Snapshot Store Lite for v3 decision outputs	Sonnet	Medium	No	Yes	Immutable per-run snapshots read/write
4	Frontend v3 held cards + drawer from v3 projection	Sonnet	High	Yes	No	At least 3 distinct action/conviction states on fixture
5	Source Pack Builder + Validator dark launch	Sonnet	High	No	Maybe	Injected hallucinations rejected
6	Committee Lite dark launch: Bull/Bear/Risk/PM	Sonnet	High	No	No	Validated outputs for sample tickers; cost cap works
7	Quality, valuation, technical modules	Sonnet	Medium-High	No	No	Modules feed axes; suppress honestly

P R	Scope	Model	Usage	UI	SQL	Stop gate
8	Forensic risk pack: filing, dilution, earnings/cash quality	Sonnet	High	No	Maybe	Known clean/flagged fixtures pass
9	Opportunity Radar dark launch, closed universe	Sonnet	High	No	Maybe	No out-of-universe candidates
10	Radar UI + committee drawer polish	Sonnet	Medium-High	Yes	No	No raw keys; radar rows explain why/risk
11	IntelToDeploy typed handoff only	Codex	Low-Med	Type only	No	Deploy math unchanged
12	Decision logs v3 + performance windows	Sonnet	Medium	Small	Yes	Provenance and horizon returns captured

Why this is better than Draft 2 roadmap

- It does not spend early budget on EDGAR or universe scans before the action model is stable.
- It puts visible product correctness before architecture breadth.
- It makes each PR testable and revertible.
- It delays Deploy integration until Intel has a stable typed contract.
- It prevents v3 from becoming another label-renaming exercise.

12. Prompt and model routing rules

Task type	Model	Chat strategy	Reason
Contract locks, tests, audits, merge gates, refactors	Codex	New unless tiny follow-up	Cheap, surgical, deterministic.
Focused multi-file implementation	Sonnet	New	Best balance for backend/frontend build PRs.
Architecture/spec/planning only	Opus	New	Use for big design decisions, not coding.
Large UI implementation	Sonnet	New with UI budget	Needs coherence but must be bounded.
Second patch after failed patch	Escalate to Sonnet Level 2	New	Stop patch loops; do full plumbing analysis.
Merge gate	Codex	New or follow-up to PR context	Review changed files only, no implementation.

Prompt wrapper requirements

- Model choice, chat strategy, usage estimate, and budget review must appear outside the copyable prompt block.
- UI prompts must include UI budget: allowed files, forbidden surfaces, visual scope, stop condition.
- Copyable prompt block contains only text to paste into Claude/Codex.
- Every non-trivial prompt includes execution principles language from docs/ai/EXECUTION_PRINCIPLES.md.
- Bug/fix prompts classify severity using docs/ai/ISSUE_SEVERITY_ROUTING.md.
- Every meaningful PR updates docs/ai/HANDOFF.md and v2/progress_log.md.

Budget gates

- Low: proceed if scope is surgical and tests are focused.
- Medium: include stop-after-PR, focused file list, no broad repo scan, and no unrelated cleanup.
- High: require phase gate, timeout budget, explicit rollback plan, and stop-after-PR. Do not combine UI + backend + SQL + LLM unless unavoidable.

13. Immediate PR 0 prompt

Use this in a new Codex chat. Do not use Sonnet for this unless Codex cannot see the needed files. This prompt intentionally rejects Draft 2 empty-stub scaffolding.

Task:

Implement Intel v3 PR 0: lock the visible Intel action contract and add regression protection against posture-label drift and HOLD-collapse. This PR must be surgical and must not implement v3 business logic yet.

Severity classification: Level 1.

Repo:

prashanthkrishnan91/claude_financetracker_pk91

Work only inside v2.

Context:

Intel v2 has repeatedly regressed into inconsistent posture labels and most recently all 34 cards showing HOLD. The v3 product contract is:

- Held tickers visible action set: BUY / HOLD / TRIM / SELL only.
- Non-held radar actions WATCH / BUY / AVOID are allowed later, but radar is not part of this PR.
- No visible posture buckets such as Add Candidate, Watchlist, Review, Risk Watch, Trim Candidate.
- Frontend must remain plain-English and beginner-friendly.
- Backend can keep transitional/internal fields, but UI must not expose posture labels.
- This PR is a contract lock and test guardrail only. Do not build Intel v3 modules yet.

Execution principles:

Before coding, state assumptions and success criteria. Keep changes simple and surgical. Every changed line must trace to this task. Fix the root cause, not a symptom. If the durable fix exceeds this scope, stop and propose the split.

Repo-local files to read FIRST:

- CLAUDE.md
- docs/ai/HANDOFF.md latest entry only
- docs/ai/PROMPT_LIBRARY.md sections: Discovery budget rule, Prompt reviewer gate
- docs/ai/EXECUTION_PRINCIPLES.md
- docs/ai/ISSUE_SEVERITY_ROUTING.md
- docs/ai/skills/implementation.md
- v2/frontend/src/app/dashboard/recommendations/page.tsx
- v2/frontend/src/components/cards/AgentInsightCard.tsx
- Existing recommendation/card tests nearest to those files

Discovery budget:

- Do not run broad find/grep/glob initially.
- Read the listed files first.
- Only inspect additional files if blocked, and state the reason.

Exact scope:

1. Lock recommendations filter tabs to exactly ALL / BUY / HOLD / TRIM / SELL.
2. Add a clear contract comment near the filter definition:
LOCKED: Intel v3 visible filter contract is ALL/BUY/HOLD/TRIM/SELL only. New visible buckets require a spec change.
3. Add or update frontend regression tests so the UI fails if any visible filter/tab/badge exposes:
Add Candidate, Watchlist, Review, Risk Watch, Trim Candidate, Strong Buy, Strong Sell, Buy More, Accumulate.
4. Ensure AgentInsightCard only renders visible held action badges BUY, HOLD, TRIM, SELL.
5. If backend sends an unknown/legacy posture label, UI must not display it. Prefer safe normalized HOLD display plus dev-safe warning only if an existing warning pattern already exists. Do not crash production UI.

6. Add a backend or frontend synthetic HOLD-collapse canary test:
 - Use synthetic fixture data, not live data.
 - Test fails if a non-degenerate recommendation set with at least one meaningful published signal collapses to all HOLD with LOW conviction.
 - The canary protects the HOLD-34 class of regression without depending on the live portfolio.
7. Do NOT create dozens of empty v3 stub files.
8. Do NOT change Deploy.
9. Do NOT change Supabase schema.
10. Do NOT implement Data Truth, Snapshot Store, Opportunity Radar, LLM Committee, or new financial modules in this PR.
11. Update docs/ai/HANDOFF.md with a short PR 0 entry.
12. Update v2/progress_log.md with a short entry.
13. README.md should not be edited unless you find a material public/setup change, which is not expected.

Tests:

- Run focused tests for changed files only.
- If a test cannot be run in the environment, state exactly why and what command should be run.

Acceptance criteria:

- Recommendations visible filters are exactly ALL / BUY / HOLD / TRIM / SELL.
- Legacy posture labels cannot appear as visible badges/tabs.
- Synthetic HOLD-collapse canary exists and fails for the all-HOLD/LOW non-degenerate case.
- No Deploy files changed.
- No SQL migration.
- No v3 module scaffolding forest.
- HANDOFF.md and v2/progress_log.md updated.
- PR opened, then stop.

Final report format:

Severity classification:

Assumptions and success criteria:

Root cause/plan:

Files changed:

Tests run:

Risks:

Supabase SQL: No

Deploy touched: No

HANDOFF.md edited:

progress_log.md edited:

README.md edited:

PR link:

14. Phase gates and acceptance criteria

Do not advance phases because a PR merged. Advance only when the gate passes.

Gate	Must be true before moving on
Gate 0: Contract stable	PR 0 merged; UI cannot show posture buckets; HOLD-collapse canary exists.
Gate 1: Kernel stable	Synthetic matrix produces BUY/HOLD/TRIM/SELL and at least two conviction bands; no Deploy touched.
Gate 2: Snapshot stable	Page-load and post-run cards render from same snapshot id; no fallback blending.
Gate 3: UI stable	Recommendations page shows clean held-card output with no raw keys and no Business Read.
Gate 4: Validator stable	Injected hallucinations are rejected: uncited claims, fake numbers, price targets, raw keys, forbidden phrases.
Gate 5: Committee stable	Bull/Bear/Risk/PM notes validate for sample tickers; deterministic action remains unchanged.
Gate 6: Radar stable	No out-of-universe tickers; top candidates are inspectable and not spammy.
Gate 7: Deploy handoff safe	Deploy math unchanged when reading typed handoff; Intel fields do not override allocation math.

Manual QA expectations

- Open Recommendations page after refresh and after page reload.
- Confirm action counts and card action badges match filters.
- Click several cards and verify drawer plain-English content.
- Verify unknown/missing data appears as missing, not as a fake negative.
- Verify no raw backend keys or posture labels leak to UI.

15. Module backlog: must, later, reject

Priority	Modules / features	Decision
Must ship early	decision_policy_lite, portfolio_governor_lite, snapshot_store_lite, source_pack_builder, committee_validator	These are the v3 spine.
Must ship for useful cards	quality/attractiveness, valuation/price, technical/timing, risk/veto summaries	Enough to differentiate Buy/Hold/Trim/Sell.
Must ship before radar UI	closed universe, screen registry, out-of-universe rejection, screen quotas	Prevents hallucinated ticker spam.
Later	full IPS Settings UI, thesis tracker, what changed, behavioral guardrails	Valuable after held-card engine is stable.
Later	EDGAR forensic depth, Form 4, 13F, customer concentration, litigation	High value but parsing/rate-limit heavy.
Later	risk stress library, factor attribution, replay, paper portfolio	Discipline layer, not needed for first v3 release.
Reject now	full 9-persona committee, paid data vendors, real-time alerts, voice CIO, auto-execution	Too expensive or unsafe before v3 kernel proves itself.

Module admission test

Before building any new module, the prompt must answer four questions. If any answer is weak, defer the module.

- Which action axis, veto, source-pack claim, UI sentence, or Deploy handoff field does this module feed?
- What inputs are required, and what happens when they are missing?
- What synthetic fixture proves the module behaves correctly?
- What exact frontend phrase will the user see, if any?

16. Risks, stop rules, and recovery plan

Primary risks

- **Architecture theater:** too many files, too many modules, no working decision output. Recovery: enforce PR 0 no-stub rule and module admission test.
- **HOLD-collapse returns:** global insufficient-data gate or blanket conviction downgrade. Recovery: canary and per-axis suppression tests.
- **LLM drift:** committee narrative disagrees with deterministic action. Recovery: schema has no action field; validator rejects action language.
- **Provider fragility:** yfinance/EDGAR/Finnhub gaps produce NaNs or stale data. Recovery: coerce_or_suppress and stale banners.

- **Deploy regression:** Intel handoff accidentally changes allocation math. Recovery: delay handoff and use Codex mechanical PR with end-to-end tests.
- **UI complexity creep:** frontend becomes giant stock report. Recovery: drawer progressive disclosure and no raw keys.

Patch exhaustion rule

- After one failed patch, reclassify severity.
- After two related patches, stop patching and run a Level 2 full plumbing analysis/fix prompt.
- If fix combines UI, backend, persistence, LLM, and tests, split unless the root cause truly requires one PR.

Recovery if a PR goes sideways

- If PR changes unrelated surfaces, request cleanup before merge.
- If tests are missing for a contract change, block merge.
- If a UI PR touches Deploy or unrelated pages, split.
- If Sonnet times out, stop with checkpoint; continue at most twice; then move to new chat.

17. Final recommendation

Build v3, but build the spine first. The goal is not the largest plan; it is the smallest architecture that permanently ends posture-label drift, HOLD-collapse, and unsourced reasoning while creating a path to radar and Deploy integration.

The correct immediate action is PR 0 with Codex in a new chat. Do not start Data Truth, EDGAR, Opportunity Radar, LLM committee, or Deploy handoff until the visible contract and canary are merged. After PR 0, bring the result back and generate PR 1: Decision Kernel Lite dark launch.

The principle to defend

A source-backed personal investment committee is achievable in this stack. The only way to get there without wasting tokens is to make each PR produce one irreversible improvement in trust: a contract, a canary, a deterministic policy, a snapshot, a validator, a committee, a radar. In that order.

End of Draft 3.